

Reverse-engineering RLC Networks With In-Circuit Measurement

by

Joshua Michael Gordonson

Submitted to the Department of Electrical Engineering and Computer Science
in partial fulfillment of the requirements for the degree of

Master of Engineering in Electrical Science and Engineering

at the

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

September 2015

© Joshua Michael Gordonson, MMXV. All rights reserved.

The author hereby grants to MIT permission to reproduce and to distribute publicly
paper and electronic copies of this thesis document in whole or in part in any
medium now known or hereafter created.

Author
Department of Electrical Engineering and Computer Science
September 08, 2015

Certified by
Gerald Jay Sussman
Panasonic Professor of Electrical Engineering
Thesis Supervisor

Accepted by
Dr. Christopher Terman
Chairman, Masters of Engineering Thesis Committee

Reverse-engineering RLC Networks With In-Circuit Measurement

by

Joshua Michael Gordonson

Submitted to the Department of Electrical Engineering and Computer Science
on September 08, 2015, in partial fulfillment of the
requirements for the degree of
Master of Engineering in Electrical Science and Engineering

Abstract

The hands-on side of electrical engineering is still taught using solderless breadboards. To lower the learning curve and improve the utility of solderless breadboards, I have designed and implemented a prototype that draws schematic diagrams of passive circuits that are built on a breadboard. The system reverse engineers circuits by means of a network sensing algorithm, which iteratively grounds and excites nodes with voltage sources, and subsequently measures the resulting currents and voltages in the network. Both a software simulation and a hardware implementation were built to test the network sensing algorithm. The software system is capable of reverse-engineering arbitrarily sized RLC networks with some caveats regarding high-q parallel RLC networks. The hardware system is able to accurately detect resistive and capacitive networks with eight nodes, though current hardware limitations significantly reduce the precision of measurement. The performance of the hardware system was analyzed and solutions to many of the measurement issues were found. A technique for surface-mount soldering breadboards to PCBs is presented in this thesis.

Thesis Supervisor: Gerald Jay Sussman
Title: Panasonic Professor of Electrical Engineering

For Jim Roberge

Thank you:

Gerald Sussman

Anne Hunter

Ed Moriarty

Lynn and Steve Gordonson

and countless friends

I could not have done this without your overwhelming support.

Contents

1	Introduction	13
1.1	History of Breadboards	13
1.2	Modern Breadboards	13
1.3	Proposed Solution	14
2	Network Sensing Algorithm	17
2.1	Network Sensing	17
2.1.1	Two Node Network	17
2.1.2	Three Node Network	18
2.1.3	N Node Network	19
2.1.4	Element Identification	20
2.1.5	Reconstructing the Network	24
3	Simulation	25
3.1	NgSpice and Netlists	25
3.2	Methods	26
3.2.1	Generate Random Netlist	26
3.2.2	Inserting Voltage Sources, Grounds and Probes	27
3.2.3	Run Simulation	27
3.2.4	Print Matrix	28
3.2.5	Write JSON	28
3.3	Executing NSA	28
3.3.1	Calculate $Z_{n }(f)$	28
3.3.2	Calculate $V_n(f)$	28
3.3.3	Calculate $Z_{nm}(f)$	29
3.3.4	Finite Difference	29
3.3.5	Element Identification	29
3.3.6	Network Reconstruction	29

4 Hardware	31
4.1 Node Voltage Reading	32
4.1.1 Onboard A/D converters	32
4.2 Signal Generator	33
4.3 Test Voltage Current Sensing	33
4.4 High-side Switches	34
4.5 Low-side Switches	34
4.6 PCB Mounted Breadboard	34
4.7 Hardware Prototypes	38
5 Firmware	43
5.1 USB	43
5.2 ADC	44
5.2.1 ADC Timers	45
5.2.2 ADC DMA	46
5.2.3 Data Reconstruction	47
5.3 DAC	48
5.3.1 DAC Wavetables	49
6 Software Control	51
6.1 Control Routine	51
6.1.1 FFT	52
6.2 Network Sensing Algorithm	52
7 Results	53
7.1 Simulation Performance	53
7.2 Hardware Performance	54
7.2.1 Element Testing	54
7.2.2 Network Testing	59
7.3 Improvements	61
7.3.1 Variable Gain Amplifiers	61
7.3.2 Variable Sense Resistor	61
7.3.3 Better Switches	62
7.3.4 Better Schematic Display	62
7.3.5 Additional Devices	62
7.3.6 Computer-Aided Debugger	62
7.4 Conclusion	63

List of Figures

2-1	Finding R_{AB} in a two node network	18
2-2	Finding R_{AB} in a three node network	18
2-3	Collapsing five node network, collapsed to find R_{AE} and R_{BC}	19
2-4	Example RLC Tank Circuit	21
2-5	Finite difference of an RLC tank where $R=6\Omega$, $L=20mH$, and $C=500nF$	22
2-6	Finite difference of an RLC tank where $R=44\Omega$, $L=1.1mH$, and $C=500nF$	23
3-1	Three node network	25
4-1	Simplified block diagram of hardware	31
4-2	ADCS7476 Equivalent Input Circuit	32
4-3	Difference Amplifier Schematic	33
4-4	MCP6L92 Open Loop Bode Plot	34
4-5	Resistor Divider Parasitic Capacitance	35
4-6	Solderless breadboard plastic face with finger spring rails removed	35
4-7	Breadboard rails hand-mounted on copper clad PCB	36
4-8	3D printed finger spring jigs	36
4-9	Top view of 36-position finger spring jig	37
4-10	Top view of 90-position finger spring jig	37
4-11	Fully loaded 36-position finger spring jig	37
4-12	Front and back of PCB mounted breadboard	38
4-13	Schematic of full hardware system	39
4-14	Layout of full hardware system	40
4-15	Top layer of full hardware system	40
4-16	Bottom layer of full hardware system	40
4-17	Top side of completed hardware system - power jack, voltage reading multiplexer, low-side switches, and 8-pin breadboard header are visible.	41

4-18	Bottom side of completed hardware system - voltage regulators, A/D converters, high-side switches, and current-sense amplifier are visible	41
4-19	8-node circuit sensing breadboard in operation	41
5-1	ADCS7476 Serial Interface Timing Diagram	44
5-2	Full ADC Driver with DMA	45
5-3	ADC Chip Select Timer	46
5-4	ADC Clock Timer	46
5-5	ADC Clock Timer Triggers DMA2	47
5-6	DAC Timer Triggers DMA1	48
7-1	Plot of measured resistance for known resistors	54
7-2	Driving point impedance of an open node from 1Hz to 100KHz	55
7-3	Impedance between two open nodes from 1Hz to 100KHz	56
7-4	Plot of measured capacitance for known capacitors	57
7-5	Impedance vs frequency for a $3.3mH$ inductor, demonstrating inductive behavior albeit noisily	58
7-6	Impedance vs frequency for a $3.3mH$ inductor in parallel with a $0.15\mu F$ capacitor, demonstrating LC behavior albeit noisily	58
7-7	Schematics of 3, 4, 5, and 6-node resistor networks rendered in the browser	59

List of Tables

5.1 <code>datas[0:16]</code>	48
--	----

Chapter 1

Introduction

1.1 History of Breadboards

The breadboard has been a staple substrate for electronic construction over the last century. At the dawn of a growing interest in amateur radio, resourceful tinkerers used planks of wood to secure and ruggedize their electrical handiwork. Conductive nodes, such as nails or copper rails, were driven into the non-conductive boards, providing anchors and contact points that were electrically isolated from the rest of the circuit. Components were soldered or wire-wrapped to the nodes, and sometimes secured by non-energized nails or screws. This construction technique provided a lot of artistic freedom in circuit construction, but was time consuming and required relatively heavy hand tools such as a hammer or drill. From a performance perspective, sensitive high-frequency circuits were infeasible due to a lack of a ground plane and long wires between components.

1.2 Modern Breadboards

The solderless breadboard is the canonical tool given to students taking introductory courses in the field. Rather than driving nodes into arbitrary locations, component leads are inserted into contact points arranged on a grid that allow rapid semi-rigid construction of circuits with no other tools. The layout of a solderless breadboard is designed to be compatible with a plethora of powerful integrated circuits - enabling complex electronic designs - but is still limited to low-frequency operation due to the parasitic capacitance between adjacent breadboard rails. Modern solderless breadboards have come a long way from their namesake wooden ancestors, but there is still room for improvement.

The intent of breadboarding is to physically realize a circuit. Often, this involves designing or using a reference schematic to guide construction, but circuit improvisation is not uncommon. A meticulous breadboarder can successfully realize a circuit without error by correctly placing compo-

nents and jumper wires - taking care not to introduce undesired 'parasitic' components. However, for the uninitiated it is difficult to justify the additional time and care required to plan and build. Inserting components and jumper wires into contact points is straightforward, but poor contacts, broken wires (inside insulation), and mis-inserting leads can plague designers for hours on end and potentially destroy components. Breadboarding is a skill that is learned over time, but small errors can lead to excessive frustration and turn students off to the field. I propose a solution to some of the issues with the solderless breadboard.

A confident linkage between the the electrical and mechanical domains is required to construct a circuit. On larger scales, the mechanical structures (eg. contacts, wires, components) that create electrical circuits are visible in plain sight. It is simple and reliable, then, to determine the circuit representation of a mechanical structure that has no hidden connections. Breadboards often obscure connections due to their construction. The regular nature of a breadboard, a grid of contact points arranged in rails, makes it difficult to keep track of which rail is connected to what circuit node. When combined with the opaque nature of most components and the poor reliability of breadboard contacts, visually verifying complex circuits on a breadboard becomes infeasible. Many of the inherent problems with breadboards stem from this open-loop nature of breadboarding, where visual cues are not enough to determine the electrical circuit from the mechanical structure. A symptom of open-loop construction techniques is a mismatch between the mental model and the physical realization of the system at hand. I seek to close this loop by designing and implementing a circuit-sensing breadboard.

1.3 Proposed Solution

A circuit-sensing breadboard is able to reverse-engineer circuits built on it by applying test voltages and currents and measuring the results. As a proof-of-concept, the proposed circuit-sensing breadboard completed for this thesis is able to reverse engineer circuits composed of resistors built on a small section of the breadboard and display a schematic of the circuit in the browser.

The circuit-sensing breadboard is composed of a hardware system that interfaces various pieces of test equipment with eight rails on a breadboard. The system is composed of a pcb-mounted breadboard, a voltage source, current meter, and voltage meter multiplexing board, a microcontroller development board, firmware to control ADC sampling, AC voltage source testing, and streaming data to a computer for analysis and display. The hardware is controlled by a circuit sensing algorithm, which determines the equivalent circuit elements between any two nodes on the breadboard. A software simulation of the hardware system verifies operation of the network sensing algorithm and enables testing for networks larger than eight nodes. Once the circuit sensing algorithm reconstructs the entire circuit it generates a schematic diagram of the circuit, closing the loop on breadboard

construction.

Chapter 2

Network Sensing Algorithm

In this chapter, a Network Sensing Algorithm that can reverse-engineer an arbitrary n-node RLC network given access to every node is described. First, it is shown how basic circuit theory is used to analyze and reverse engineer n-node resistive circuits. These methods are then generalized to reverse-engineer n-node impedance networks down to the individual impedance branch circuits, constructed from resistors, capacitors, and inductors. The impedance networks are then further decomposed into their constituent element by method of frequency-domain analysis, at which point the original network is fully represented and can be reconstructed.

2.1 Network Sensing

Given access to the set of nodes in an electrical network, the objective of a network sensing algorithm is to determine the set of branches and the elements that compose them. In order to illustrate how a network sensing algorithm operates, it is necessary to begin with a simple example and then build in complexity.

2.1.1 Two Node Network

Take the example of a network with two nodes below:

In the case of a resistive network with two nodes, there is only one possible branch in the network and thus one possible element to characterize. From Section 2.1, the resistance between two nodes equal to the voltage across the nodes divided by the current through the nodes when power is applied. Here, the resistance R_{AB} is found by placing a test voltage V_t across the nodes and measuring the resulting current, I_t , then taking the ratio $\frac{V_t}{I_t}$. This measurement is called the driving point impedance¹. If there are no circuit elements between the two nodes, the driving point

¹To measure a driving point impedance, a test voltage is applied between the node of interest and ground,

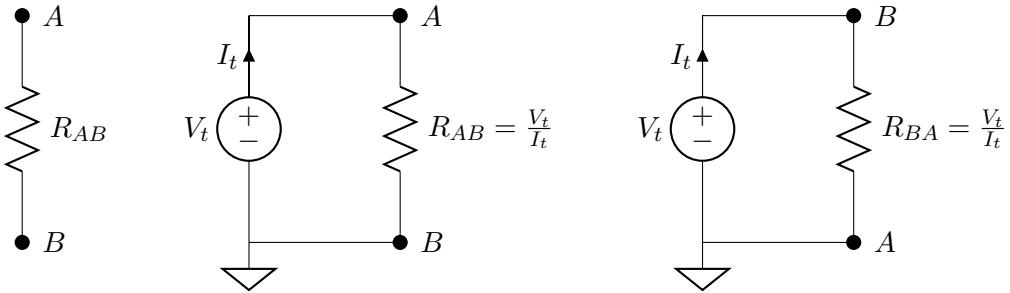


Figure 2-1: Finding R_{AB} in a two node network

impedance test will find zero current in the test voltage source, resulting in an infinite resistance between two nodes. An infinite resistance between two nodes in a circuit indicates that there are no elements in that branch of the network. Thus, the network can subsequently be simplified by removing that branch from the network.

2.1.2 Three Node Network

A resistive network with two nodes is simple, but provides an introduction to the methodology used in analyzing larger networks. In the case of a resistive network with three nodes, it is insufficient to utilize driving point impedance measurements alone, because each node has more than one path to any other node. The ability to ground arbitrary nodes of a network greatly simplifies the process of analyzing and reverse-engineering a network. The network-sensing algorithm relies on the ability to shrink the effective network by conglomerating nodes into the ground node.

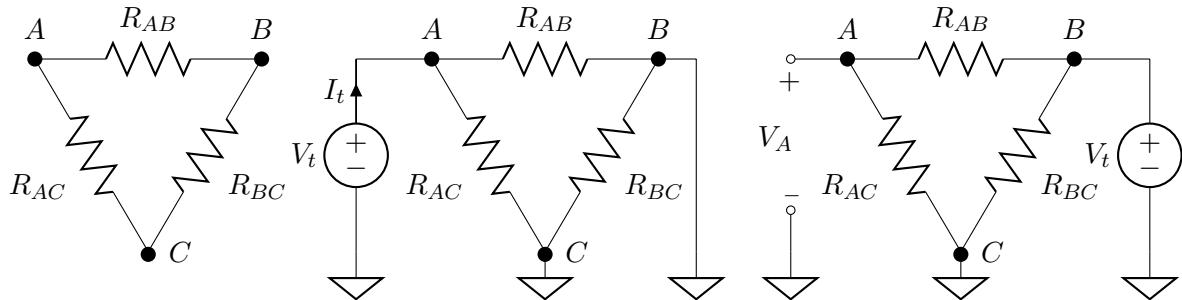


Figure 2-2: Finding R_{AB} in a three node network

Consider a resistive network with three nodes: A, B, and C. In order to determine the resistance in branch AB, the driving point impedance at node A is measured with nodes B and C grounded. This provides the resistance of the parallel combination of the branches with an endpoint at node

and the resulting current in the voltage source is measured. The driving point impedance is then computed by dividing the test voltage by the resulting current.

A:²

$$R_{A||} = R_{AB} \parallel R_{AC}$$

Next, a test voltage source is applied to node B, node C is grounded, and the voltage at node A, V_A , is observed.

$$V_A = V_t \frac{R_{AC}}{R_{AB} + R_{AC}} = V_t \frac{R_{AB} \parallel R_{AC}}{R_{AB}} = V_t \frac{R_{A||}}{R_{AB}}$$

The branch resistance of interest, R_{AB} is calculated using the known quantities V_t , $V_{A||}$, and V_A .

$$V_t \frac{R_{A||}}{V_A} = V_t \frac{R_{AB} \parallel R_{AC}}{V_t \frac{R_{AC}}{R_{AB} + R_{AC}}} = \frac{\frac{R_{AB} R_{AC}}{R_{AB} + R_{AC}}}{\frac{R_{AC}}{R_{AB} + R_{AC}}} = R_{AB}$$

This procedure is repeated for the remaining branches to determine the entire network.

2.1.3 N Node Network

A resistive network with any number of nodes can be reduced to a resistive network with three nodes by grounding the nodes that are not of interest. The resulting network does not modify the branch of interest, but connects the remaining branches attached to the nodes of interest in parallel. This collapses the network into a three node network or three branch equivalent circuit. Consider

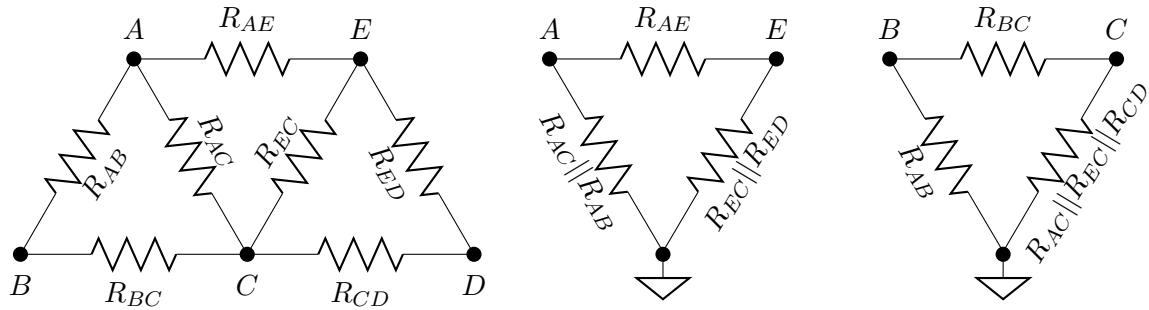


Figure 2-3: Collapsing five node network, collapsed to find R_{AE} and R_{BC}

a resistive network with five nodes: A-E. To determine the resistance between nodes A and E, the network is reduced to a three-node network by connecting all nodes except nodes A and E to ground. The three resistances of the branches that remain are R_{AE} , $R_{AB} \parallel R_{AC}$, and $R_{EC} \parallel R_{ED}$. The reduced network is then solved using the three node network method, and this procedure is repeated for all branches in the network.

² $R_1 \parallel R_2 = \frac{R_1 R_2}{R_1 + R_2} = \frac{1}{\frac{1}{R_1} + \frac{1}{R_2}}$

The above thought experiment satisfies the problem at hand intuitively, but it can also be shown mathematically.

To find R_{AB} of an n-node complete network, first the driving point impedance at node A, $R_{A_{||}}$, is measured with all other nodes grounded.

$$\begin{aligned} R_{A_{||}} &= R_{AB} \parallel R_{AC} \parallel R_{AD} \parallel \dots = \frac{1}{\frac{1}{R_{AB}} + \frac{1}{R_{AC}} + \frac{1}{R_{AD}} + \dots} \\ &= \frac{1}{\frac{1}{R_{AB}} + \frac{1}{(\frac{1}{R_{AC}} + \frac{1}{R_{AD}} + \dots)}} \\ &= R_{AB} \parallel (R_{AC} \parallel R_{AD} \parallel \dots) \end{aligned}$$

To find V_A from node B, the voltage at node A is measured when a test source is placed at node B with all other nodes grounded.

$$V_A = V_t \frac{(R_{AC} \parallel R_{AD} \parallel \dots)}{R_{AB} + (R_{AC} \parallel R_{AD} \parallel \dots)}$$

Dividing $R_{A_{||}}$ by V_A and multiplying by V_t :

$$V_t \frac{R_{A_{||}}}{V_A} = V_t \frac{\frac{R_{AB}(R_{AC} \parallel R_{AD} \parallel \dots)}{R_{AB} + (R_{AC} \parallel R_{AD} \parallel \dots)}}{\frac{(R_{AC} \parallel R_{AD} \parallel \dots)}{R_{AB} + (R_{AC} \parallel R_{AD} \parallel \dots)}} = R_{AB}$$

Finding a resistive element in an n-node network requires two measurements. The number of branches in a complete graph with n nodes is equal to $\frac{n(n-1)}{2}$. The total number of measurements required to reverse-engineer an n-node circuit is $n(n - 1)$.

2.1.4 Element Identification

Networks composed of elements with complex impedance can be analyzed with the same algorithm. By replacing the test DC voltage sources with AC voltage sources, the reactance of capacitors and inductors can be measured in addition to the resistance of resistors.

Parallel RLC Branches

To determine if there are multiple element types [in parallel] between two nodes, the impedance is measured over a range of frequencies, and the resulting changes in impedance are measured.

The impedance of a parallel RLC 'tank' circuit can be characterized and analyzed over all frequencies.

$$Z_R = R \quad Z_L = j\omega L \quad Z_C = \frac{1}{j\omega C}$$

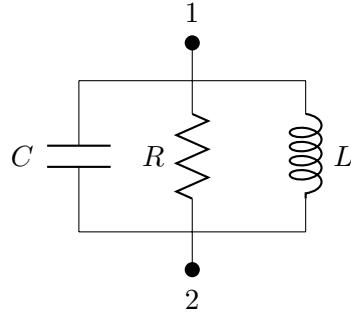
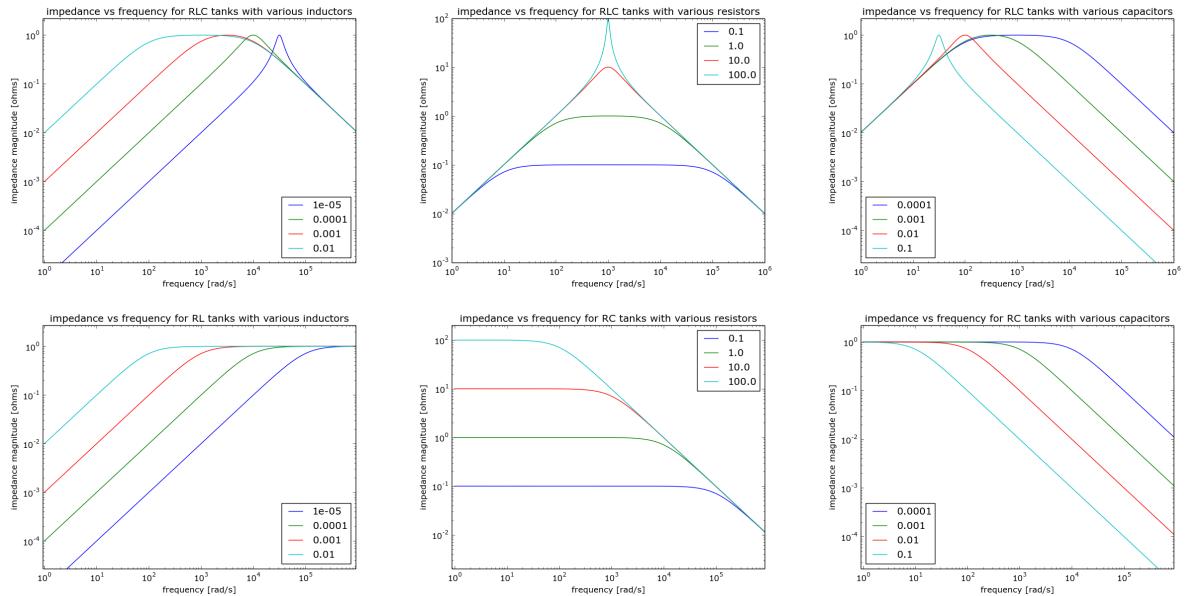


Figure 2-4: Example RLC Tank Circuit

$$Z_{RLC} = Z_C || Z_R || Z_L = \frac{1}{j\omega C + \frac{1}{R} + \frac{1}{j\omega L}} = \frac{j\omega RL}{-\omega^2 RLC + j\omega L + R}$$

Z_{RLC} has a zero at $\frac{1}{RL}$ and two poles at $\frac{-L \pm \sqrt{L^2 + 4R^2 LC}}{-2RLC}$



The plots above illustrate various RLC tank circuits where one component value changes over several orders of magnitude. In the top-center and bottom-center plots, a decrease in resistance drops the maximum impedance over all frequencies. In the top-left and bottom-left plots, a decrease in inductance drops the asymptotic impedance at low frequencies. In the top-right and bottom-right plots, a decrease in capacitance drops the asymptotic impedance at high frequencies. As illustrated above, there are three regimes on the impedance vs frequency plot that divide the operation of an RLC tank into three elements:

When the slope of $|Z|$ vs. f is +1, the tank behaves like an inductor with inductance

$$L = |Z|/(j\omega)$$

When the slope of $|Z|$ vs. f is -1, the tank behaves like a capacitor with capacitance

$$C = j\omega|Z|$$

When the slope of $|Z|$ vs. f is 0, the tank behaves like a resistor with resistance

$$R = |Z|$$

Finite Difference Stencil

To reliably determine the regions of effective resistance, inductance, and capacitance, the slope of the $\log(|Z_{nm}|)$ vs. $\log(\omega)$ is computed via the midpoint method. The midpoint finite difference stencil takes the two points on either side of the point of interest and computes the slope between them.

$$Z'[n] = \frac{1}{2}(Z[n+1] - Z[n-1])$$

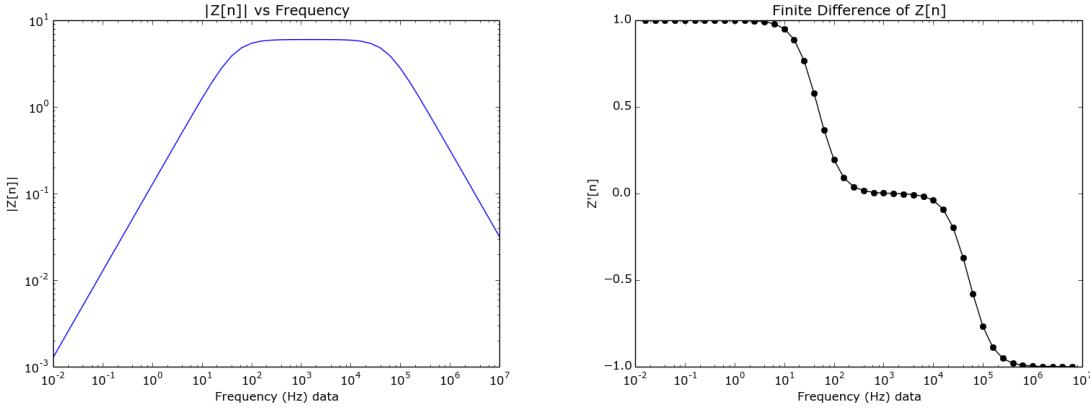


Figure 2-5: Finite difference of an RLC tank where $R=6\Omega$, $L=20mH$, and $C=500nF$

In figure 2-5 the distinct regions of inductance, resistance, and capacitance are clearly visible and easily found by setting thresholds for the finite difference. The thresholds on the region of inductance are between 1.1 and 0.9, the thresholds on the region of resistance are between 0.1 and -0.1, and the thresholds on the region of capacitance are between -0.9 and -1.1.

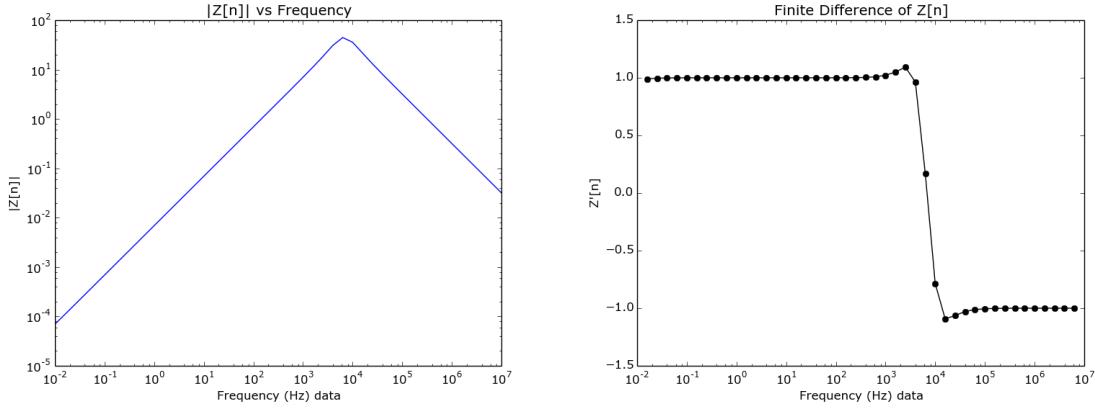


Figure 2-6: Finite difference of an RLC tank where $R=44\Omega$, $L=1.1\text{mH}$, and $C=500\text{nF}$

In figure ??, the resistance is not represented in the finite difference of $Z[n]$. The parallel resistance of an underdamped RLC tank is the maximum recorded impedance across all frequencies. In a simulation, using the maximum recorded impedance is a valid solution because measurement accuracy has a large (64-bit floating point) dynamic range, producing results that are accurate of ideal systems. In practice, spurious noise and lack of dynamic range over measurement data can cause poor results. Occasionally these poor results include large spikes of inaccurate impedance readings. In this case, a practical approach would be to use the finite difference stencil where possible and when an LC network is discovered, sample at the resonant frequency to find R . In the case of figure 2-6, the sampling frequency would be 17KHz.

$$f_0 = \frac{1}{\sqrt{2\pi LC}} \approx 17000\text{Hz}$$

Component Value Calculation

Once the regions of the impedance plot are thresholded, the sampling frequency domains corresponding to those regions are assigned to an element type [$\omega_r[n]$ for R, $\omega_l[n]$ for L, $\omega_c[n]$ for C]. Then for each sampling frequency the component values are computed and averaged to filter noisy data.

$$R[\omega_r] = Z[\omega_r] \quad L[\omega_l] = \frac{Z[\omega_l]}{\omega_l} \quad C[\omega_c] = \frac{1}{\omega_c Z[\omega_c]}$$

$$R = \frac{1}{rnum} \sum_{\omega_r} Z[\omega_r] \quad L = \frac{1}{lnum} \sum_{\omega_l} \frac{Z[\omega_l]}{\omega_l} \quad C = \frac{1}{cnum} \sum_{\omega_c} \frac{1}{\omega_c Z[\omega_c]}$$

Where $rnum$, $lnum$, and $cnum$ are the number of sample frequencies in the R, L, and C regions, respectively.

2.1.5 Reconstructing the Network

With a record of all of the elements and their connections in the network, the network can be reconstructed and a schematic of the original circuit can be drawn.

Chapter 3

Simulation

In this chapter, a simulation of the circuit sensing breadboard is described. A simulated circuit sensing breadboard is useful for rapid iteration of hardware topologies and validation of network sensing algorithms. The simulator is also not limited by the hardware complexity, so simulating circuits with more than eight nodes and elements other than resistors, capacitors, and inductors is possible. The simulator generates random networks of resistors, inductors, and capacitors, and proceeds to analyze, reconstruct, and draw a schematic of the network using the network sensing algorithm.

3.1 NgSpice and Netlists

The open source software package NgSpice was used to simulate the networks and test circuits applied to the network. NGspice operates on text files called netlists, where each component in the network is specified on a single line. An example netlist is shown below.

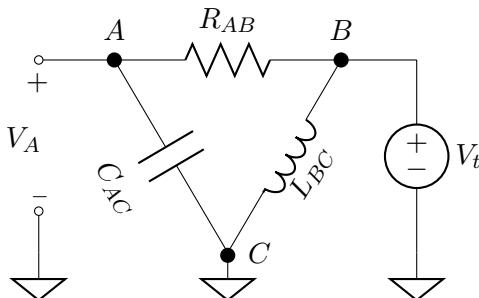


Figure 3-1: Three node network

```
threeNodeNetlist
Vt      0 B 1
```

```

Rab A B 10k
Cac A C 1e-6
Lbc B C 1e-3
Vcg 0 C 0
.control
    op
    print(v(A))
    .endc
.end

```

The first letter of each element line designates the element to simulate:

V start stop value → Voltage Source between **start** and **stop** nodes, **value** Volts [V]

R start stop value → Resistor between **start** and **stop** nodes, **value** Ohms [Ω]

L start stop value → Inductor between **start** and **stop** nodes, **value** Henries [H]

C start stop value → Capacitor between **start** and **stop** nodes, **value** Farads [F]

Node 0 is always designated as ground, and all simulations require a ground node.

The control commands are as follows:

op → Operating Point Simulation **print()** → Print the relevant data passed as an argument **v(N)**

→ The voltage at node N **i(E)** → The current through element E

3.2 Methods

The NSA simulator was written in python and uses the methods below.

3.2.1 Generate Random Netlist

```
writeRandomNet(netlist,num,elements):
```

Generates **num**-node random graph with no self-linking nodes (symmetric matrix with zeros on the diagonal) for each **[elements]** type (R,L,C). Subsequently assigns random values between two realistic limits for each element and writes the network to netlist **netlist**. 1-1k ohms, 10nF-10uF, 100uH-100mH.

When writing the capacitive and inductive elements, care must be taken to prevent a DC operating point simulation from failing. The infinite resistance across a capacitor and zero resistance across an inductor are responsible for DC operating point simulation failure, and can be fixed by including a small resistor in series with inductors and a large resistor in parallel with capacitors.

$L0\ 1\ 2\ 1mH \rightarrow L0\ 1\ t10\ 1mH\ R10$ $t10\ 2\ 1e-5$	$C0\ 1\ 2\ 1e-6F \rightarrow C0\ 1\ 2\ 1e-6F\ Rc0$ $1\ 2\ 1e8$
--	---

3.2.2 Inserting Voltage Sources, Grounds and Probes

```
insertProbe2(target, nodes, groundNodes, probes, source='DC'):
```

Inserts a 1V voltage source from ground to each node in `[nodes]`, grounds each node in `[groundNodes]`, and adds a voltage print statement for each node in `[probes]`. By default, the voltage sources are written as DC sources, but if 'AC' is passed into the last argument the sources are written as AC sources and the AC control statement is added.

```
AC dec 5 10m 10meg
```

Which runs a small-signal AC simulation and returns the amplitudes of the resulting voltage and current waveforms, five sample points per decade from 10mHz to 10MHz.

3.2.3 Run Simulation

```
def runSim(target, results, source='DC'):
```

Makes a system call to NgSpice in batch mode with netlist `target` and outputs the result to text file `results`. The last argument indicates how to parse the resulting data, as NgSpice returns DC data in the following format:

```
No. of Data Rows : 1
i(v) = -1.18295e-01
v(5) = 2.532846e-07}
```

and AC data is returned in this format:

```
No. of Data Rows : 6
mynetlist
AC Analysis Sun Aug 30 18:35:07 2015
```

Index	frequency	i(v)
0	1.000000e+00	-1.72598e+00, 3.978810e+02
1	1.000000e+01	-1.58689e-01, 3.978827e+01
2	1.000000e+02	-1.43015e-01, 3.974287e+00
3	1.000000e+03	-1.42859e-01, 3.520201e-01
4	1.000000e+04	-1.42857e-01, -4.18884e-01
5	1.000000e+05	-1.42857e-01, -4.58275e+00

```
mynetlist
AC Analysis Sun Aug 30 18:35:07 2015
```

Index	frequency	v(3)
-------	-----------	------

0	1.000000e+00	1.000000e+00,	0.000000e+00
1	1.000000e+01	1.000000e+00,	0.000000e+00
2	1.000000e+02	1.000000e+00,	0.000000e+00
3	1.000000e+03	1.000000e+00,	0.000000e+00
4	1.000000e+04	1.000000e+00,	0.000000e+00
5	1.000000e+05	1.000000e+00,	0.000000e+00

DC data is returned as a dictionary with indices for the voltage source current and each node of interest:

```
[‘i’:i(V), ‘1’:v(1), ‘2’:v(2), ...]
```

AC data is returned as a dictionary of lists with indices for the voltage source current and each node of interest:

```
[‘i’:[i(Vf1), i(Vf2), ..., i(Vfn)], ‘1’:[v(1f1), v(1f2), ..., v(1fn)], ...]
```

3.2.4 Print Matrix

```
def printMatrix(m): Prints matrix m in nice command-line output. It can handle both 2-dimensional and 3-dimensional matrices, for n-by-n matrices with lists of impedances over many frequencies.
```

3.2.5 Write JSON

```
def writeJSON(name, network, group, elem): Writes a JSON file named name.json of the n-by-n symmetric matrix network. Three separate JSON files are written, one for each element [resistor.json, capacitor.json, and inductor.json]. The group parameter is used when drawing multiple schematics on one page. The elem parameter is used to specify what type of element the network represents.
```

3.3 Executing NSA

3.3.1 Calculate $Z_{n||}(f)$

$Z_{n||}(f)$ is found by grounding all nodes except for node n and adding a voltage source to that node, then taking the ratio of the amplitudes of the resulting current into the node of interest and the voltage source.

3.3.2 Calculate $V_n(f)$

$V_n(f)$ is found by grounding all nodes except for nodes n and m, adding a voltage source to node m, and measuring the amplitude of the voltage at node n.

3.3.3 Calculate $Z_{nm}(f)$

$Z_{nm}(f)$ is calculated by the ratio of $Z_{n||}(f)$ to $V_n(f)$ scaled by V_t . In the case of this simulation, V_t is one.

3.3.4 Finite Difference

Z_{nm} is passed through the finite-differencer, returning a list $Z'_{nm}[f]$. $Z'_{nm}[f]$ is thresholded to find the frequencies for which $Z'_{nm}[f] \approx -1$, $Z'_{nm}[f] \approx 0$, and $Z'_{nm}[f] \approx 1$.

3.3.5 Element Identification

A list of inductances is calculated from $Z_{nm}[f]$ on all frequencies for which $Z'_{nm}[f] \approx -1$, $[f_L]$.

$$L[f_L] = \frac{|Z_{nm}[f_l]|}{(f_l)}$$

A list of resistances is calculated from $Z_{nm}[f]$ on all frequencies for which $Z'_{nm}[f] \approx 0$, $[f_R]$.

$$R[f_R] = |Z_{nm}[f_R]|$$

A list of capacitances is characterized from $Z_{nm}[f]$ on all frequencies for which $Z'_{nm}[f] \approx -1$, $[f_C]$.

$$C[f_C] = \frac{1}{|Z_{nm}[f_C]|(f_C)}$$

The element lists are averaged to filter out inaccurate data.

3.3.6 Network Reconstruction

A new matrix is written for each element type, then the matrices are written to JSON and processed by javascript running D3 to display the final schematic.

Chapter 4

Hardware

In this chapter, hardware system design is covered. The hardware system is decomposed into its subsystems, and the design choices, specifications, and operation of those subsystems are explained. The hardware system is composed of a set of eight breadboard rails connected to a node voltage reading system, a test voltage source system, a grounding system, and a 32-bit microcontroller development board. The node voltage reading system is composed of a voltage multiplexer and an external A/D converter. The test voltage source system is composed of an internal D/A converter, an external buffer, a current sensing amplifier, an external A/D converter, and eight high-side analog switches. The grounding system is composed of eight N-channel MOSFETs, handled by the microcontroller. A simplified block diagram is below:

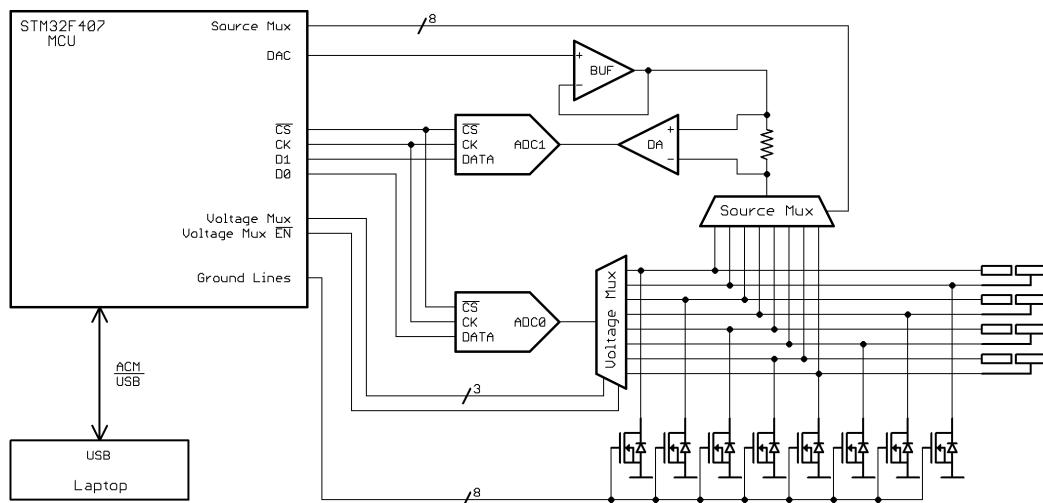


Figure 4-1: Simplified block diagram of hardware

4.1 Node Voltage Reading

An ADCS7476 12-bit A/D converter is used to measure the voltages at each node. The ADCS7476 can sample up to 1MSPS with ± 1 LSB of total unadjusted error from -40°C to 85°C and $< \pm 0.2$ LSB of error at 25°C . No bits are wasted in the ADCS7476 A/D converter. Additionally, the input circuitry to the ADC is a 100Ω resistor in series with a 26pF capacitor. This places a pole at 61MHz, causing a .03% error at 100KHz and .3% error at 1MHz. The additional performance is well worth the additional cost of \$1.56375 in quantities of 1Ku.

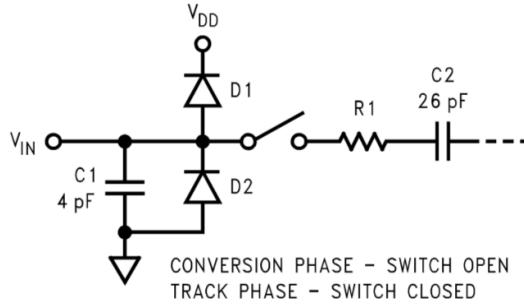


Figure 4-2: ADCS7476 Equivalent Input Circuit

The ADC is connected to the common pin of a CD4051 1:8 analog multiplexer. The multiplexer is connected to eight breadboard rails, which allows the ADC to measure the voltage on any of the eight rails, one rail at a time.

The CD4051 has about 200Ω of series resistance and 30pF of output capacitance when its supply voltage is 12V, which places an additional pole at 26MHz, again well above the Nyquist frequency.

So far, the voltage-reading signal chain has two poles - one at 61MHz and one at 26MHz.

4.1.1 Onboard A/D converters

Although the STM32F407 has three onboard 12-bit A/D converters, their specifications are lacking. Each is able to sample at 2MSPS, and it's possible to interleave them to attain a sampling rate of 6MSPS or higher if you're willing to throw away bits. The total unadjusted error (offset error, gain error, differential linearity error, and integral linearity error) is between ± 2 LSB and ± 5 LSB. With a minimum of ± 2 LSB's of error, the lowest significant two bits in the 12-bit A/D are virtually useless. Another specification to consider is the input circuitry to the ADC. The input circuitry to the ADC while it's in tracking mode looks like a $6\text{K}\Omega$ resistor charging a 4pF capacitor. This puts a pole at 6.6MHz, causing .3% error in measurement at 100KHz, 3% error in measurement at 1MHz, and 7% at 3MHz. When sampling at the maximum sample rate of 6MSPS, the ADC's input network begins to introduce significant error. Granted, it's likely that there will be an alternative

bandwidth bottleneck, this is still a metric of concern. [DM00037051.pdf, pages 133-134]

4.2 Signal Generator

The STM32F407 has an onboard D/A converter that is good enough to use as a signal generator. The onboard DAC is configured to output a cosine wavetable with a DC offset, as described in the next chapter.

4.3 Test Voltage Current Sensing

The DAC output is buffered by half of an MCP6L92 10MHz op-amp. The onboard DAC can be configured with an optional onboard buffer, but the buffer limits the DAC output range from 0.2V to Vdd-0.2V. Without the onboard buffer, the DAC output impedance is $15\text{k}\Omega$. When configured as a voltage buffer, the MCP6L92 has an input impedance larger than $1\text{G}\Omega$, resulting in no signal attenuation.

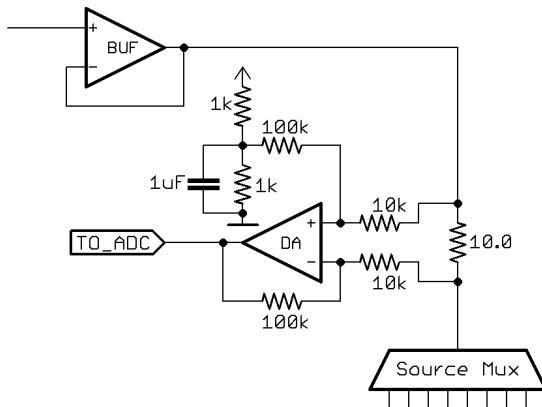


Figure 4-3: Difference Amplifier Schematic

The buffer sources current through a 10.0Ω , 1% sense resistor, which can be switched onto any of the breadboard rails through an array of eight high-side switches. The voltage across the sense resistor is measured by the other half of the MCP6L92 op-amp configured as a simple difference amplifier. Using two $10k\Omega$ resistors and two $100k\Omega$ resistors, the difference amplifier is configured with a gain of 10.

According to the MCP6L92 datasheet, the difference amplifier should have a -3dB bandwidth of 1MHz and plenty of phase margin driving the 100Ω - 26pF input impedance to the current sense ADC.

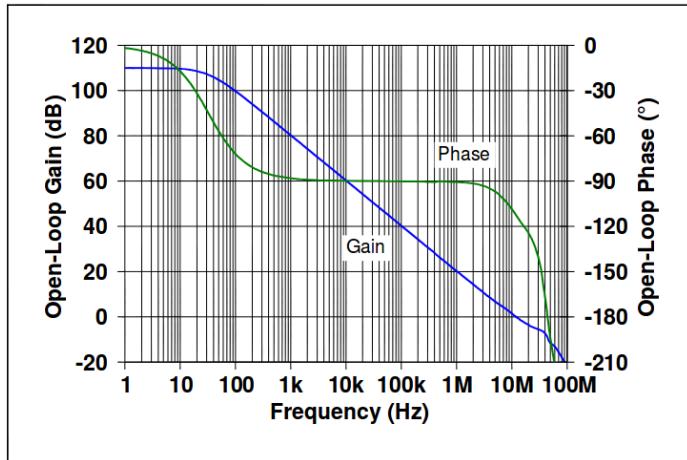


Figure 4-4: MCP6L92 Open Loop Bode Plot

4.4 High-side Switches

The high-side switches were selected for high-voltage operation, so that any rail of the breadboard could swing between 0 and 30V and there wouldn't be a problem. The selected switches were Vishay DG468 normally open analog switches. They have 9Ω resistance and 76pF of capacitance in the on-state, 1nA of leakage current and 30pF of capacitance in the off-state. The high and low side switches are the main limits of bandwidth due to their high amounts of input capacitance in both on and off states.

4.5 Low-side Switches

The low-side switches have a low logic-level threshold voltage, low drain-source on-resistance, can handle up to 30V, and are low-cost. The IRLML2803 N-FETs have an $R_{DS_{ON}}$ of about 1Ω with a V_{GS} of 3.3V. With 10V V_{GS} , $R_{DS_{ON}}$ drops to $250\text{m}\Omega$. The downside of these FETs is the high C_{DS} that comes from a wide transistor. C_{DS} is on the order of 60pF for each transistor. When combined with the additional capacitances mentioned above, each breadboard rail has a total of 140pF to ground. This has a significant impact on the maximum usable frequency for even moderate impedances on the breadboard. For example, consider a $10\text{k}\Omega$ - $10\text{k}\Omega$ resistor divider. At 100kHz, the 140 pF capacitance on each rail has $11\text{k}\Omega$ of impedance.

4.6 PCB Mounted Breadboard

Modern solderless breadboards are composed of three components: a plastic face, metal finger springs (rails), and a foam backing. The finger springs are designed to fit snugly in the plastic face, which

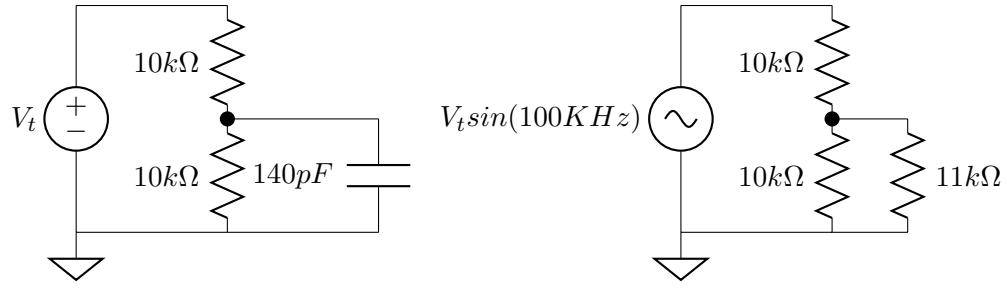


Figure 4-5: Resistor Divider Parasitic Capacitance

provides structure for the entire board. The foam backing keeps the finger springs from being pushed out the back of the plastic face..

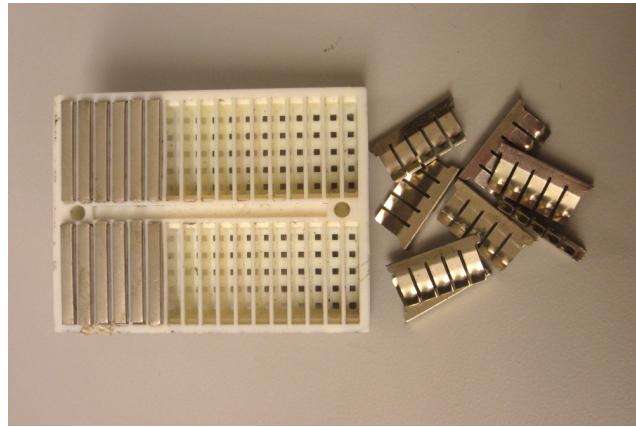


Figure 4-6: Solderless breadboard plastic face with finger spring rails removed

A breadboard can be mounted to a PCB by removing the foam backing and surface mount soldering the finger springs to a PCB, then pressing the plastic face onto the mounted rails. Initial tests of this technique involved small numbers of finger springs hand-soldered one-by-one to a PCB, using nail polish as soldermask.

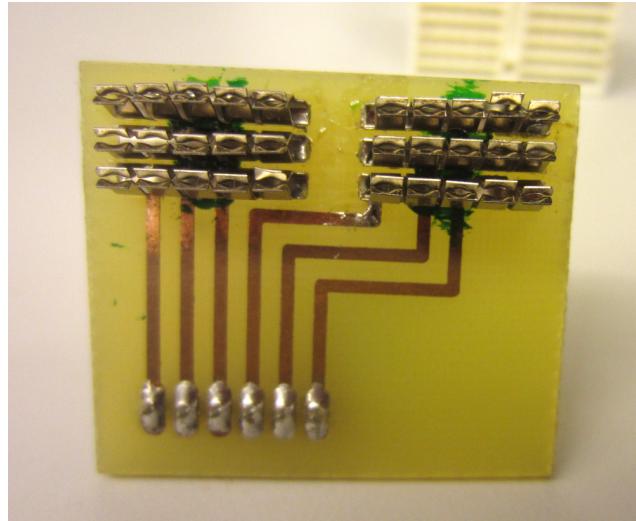


Figure 4-7: Breadboard rails hand-mounted on copper clad PCB

The results of these tests were promising, but alignment proved to be an issue. To mitigate this, a finger spring jig was 3d-printed to loosely hold each finger spring in precise alignment for PCB mounting.

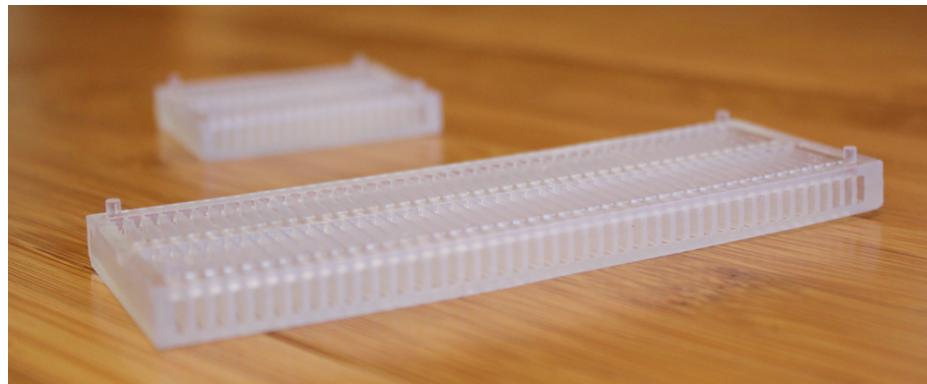


Figure 4-8: 3D printed finger spring jigs

The finger spring jig was designed in openSCAD and has four alignment pegs to mate with alignment holes on a PCB. To assemble a PCB mounted breadboard, the finger spring rails are removed from the breadboard and inserted into the finger spring jig.

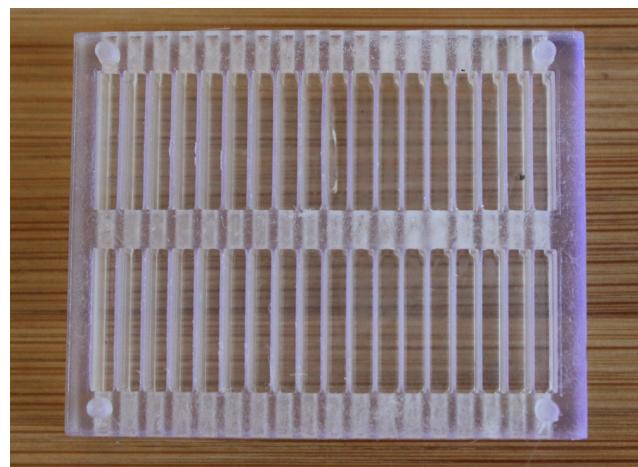


Figure 4-9: Top view of 36-position finger spring jig

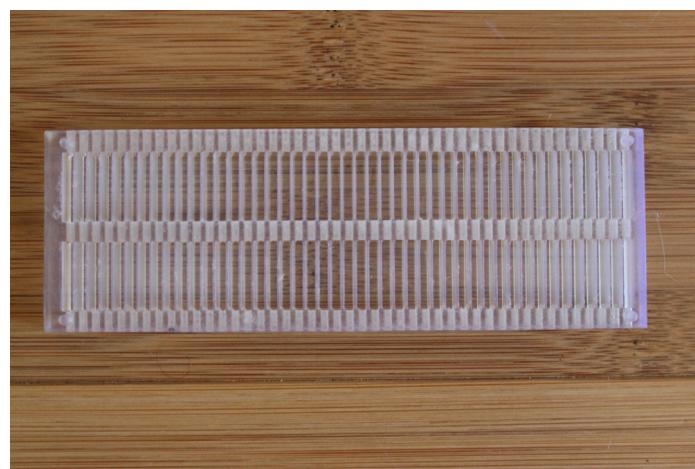


Figure 4-10: Top view of 90-position finger spring jig

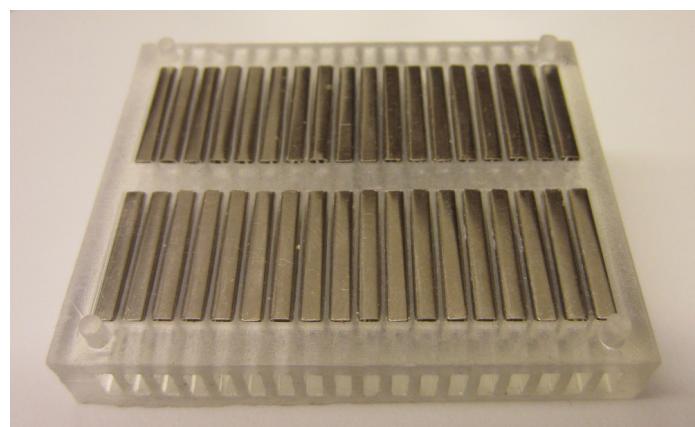


Figure 4-11: Fully loaded 36-position finger spring jig

Once all of the rails are placed, a dot of super glue is dispensed in the center of each finger spring, between two dots of solderpaste. The rails are then adhered to the PCB by pressing the PCB upside-down onto the finger spring jig while the alignment pegs are properly mated with the alignment holes in the PCB. This ensures that the array of finger spring rails are aligned with the pads on the PCB. After about ten seconds the rails adhere to the PCB and the PCB can be carefully separated from the finger spring jig. Finally the assembly is reflowed with hot air or in a reflow oven, and the plastic faceplate can be pressed on.

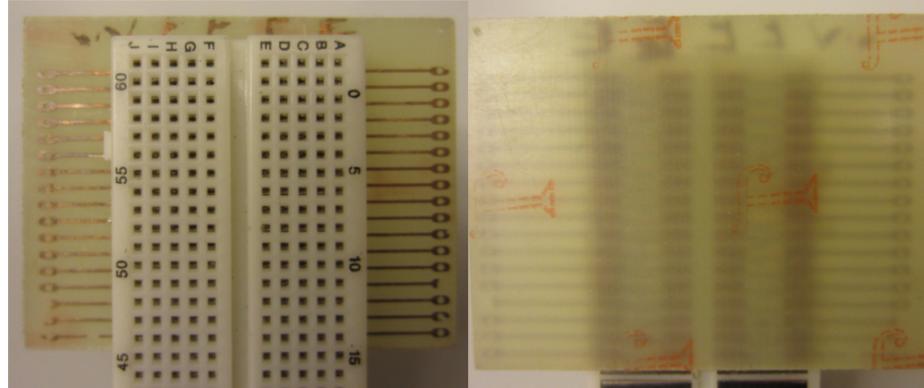


Figure 4-12: Front and back of PCB mounted breadboard

4.7 Hardware Prototypes

A number of prototype hardware subsystems were built to test various parts of the RLC network identifying breadboard. All of these systems were designed and built by hand at the MIT Electronics Research Society [MITERS] facilities.

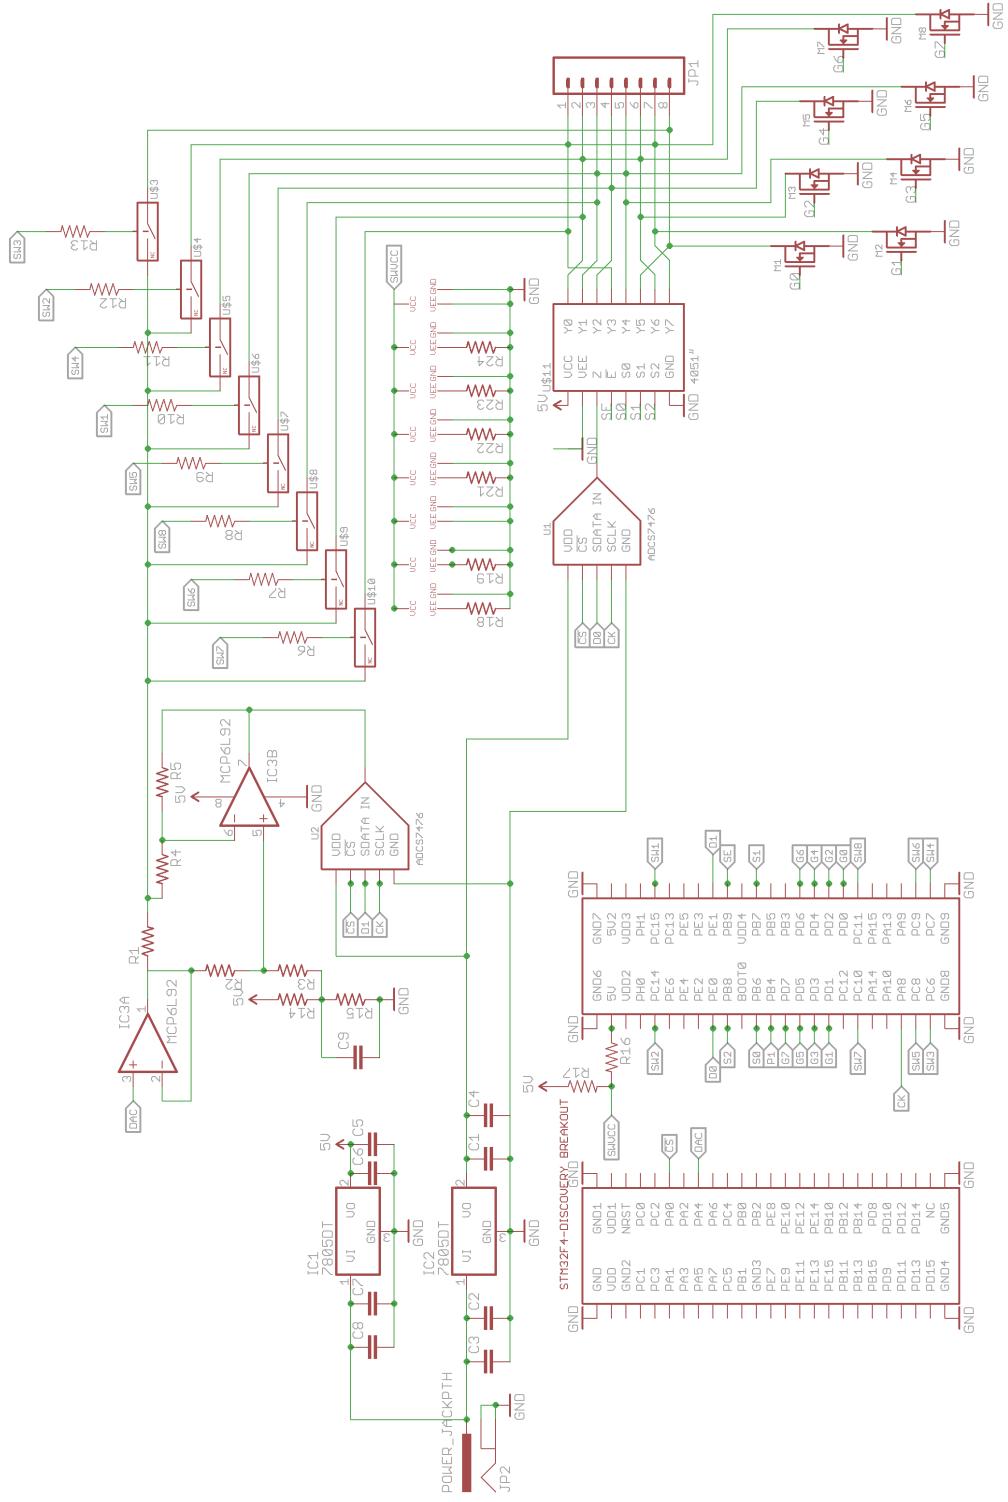


Figure 4-13: Schematic of full hardware system

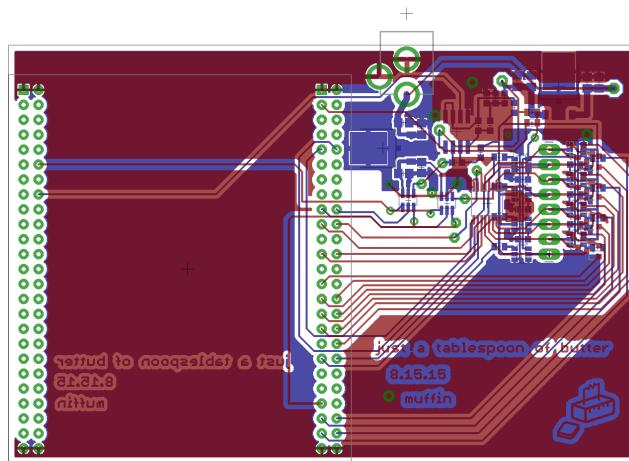


Figure 4-14: Layout of full hardware system

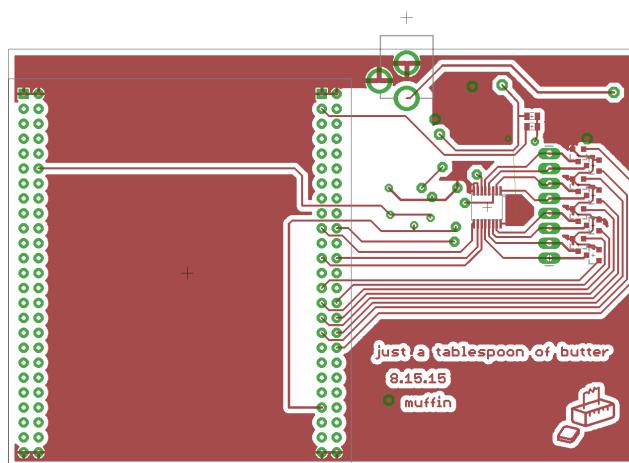


Figure 4-15: Top layer of full hardware system

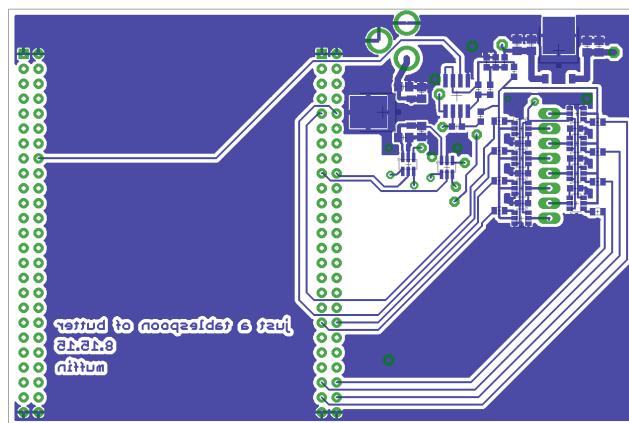


Figure 4-16: Bottom layer of full hardware system

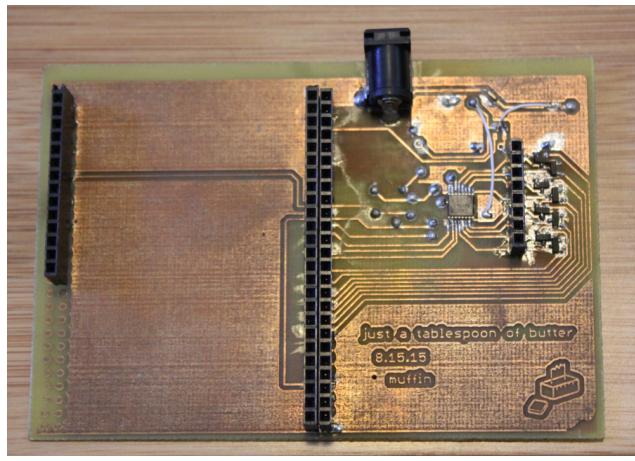


Figure 4-17: Top side of completed hardware system - power jack, voltage reading multiplexer, low-side switches, and 8-pin breadboard header are visible.

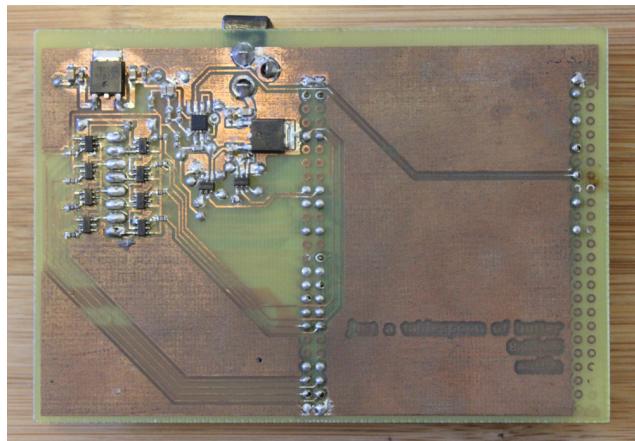


Figure 4-18: Bottom side of completed hardware system - voltage regulators, A/D converters, high-side switches, and current-sense amplifier are visible

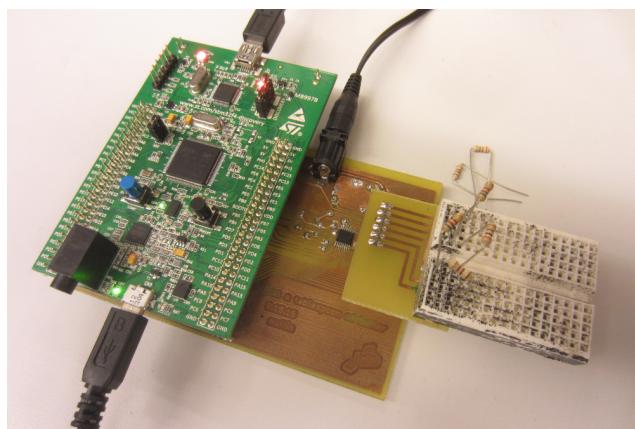


Figure 4-19: 8-node circuit sensing breadboard in operation

Chapter 5

Firmware

In this chapter, firmware configuration and operation is covered. The firmware is structured in three sections - USB data reception and transmission, DAC operation, and ADC operation. The USB subsystem receives commands from a host computer, specifying which nodes to ground, which nodes to connect to the voltage source, which nodes to probe, and what sample frequency to use. On a 'sample' command, the USB subsystem triggers both the DAC and ADC subsystems to operate synchronously, accumulating AC test samples in the memory buffer. When the buffer is full, the ADC subsystem sends the recorded data back to the host computer through the USB subsystem and waits for the next command.

5.1 USB

Libopencm3, licensed under the GNU LGPL v3, is used extensively in this firmware. Of particular utility, the USB subsystem relies entirely on example USB CDC ACM firmware included in the libopencm3 library. USB CDC (Communications Device Class) ACM (Abstract Control Module) is a USB protocol that is used to create virtual serial ports over USB. The firmware components that implement USB-CDC-ACM are a polling routine that is the only code running in the main loop and two callbacks - one for receiving data and one for transmitting data. The callback used to transmit data is called after the ADC subsystem has finished sampling and sends all of the collected data over USB. The callback used to receive data is called when the laptop sends commands over USB.

Command List

Commands are sent in sequences of characters over USBACM. The command list is as follows:

`rcvBuf=['g',nodes]` grounds all of the nodes represented by `nodes`. In this command, `nodes`

is an 8-bit binary representation of which nodes to ground. For example, `['g', 0x03]` grounds nodes 1 and 2.

`rcvBuf=['w',node]` connects the voltage source to `node` through a high-side switch. `node` is a decimal representation of which node to connect to the voltage source. For example, `['w',0x03]` connects the voltage source to node 3.

`rcvBuf=['m',node]` configures the voltage multiplexer to connect `node` to the A/D converter. `node` is a decimal representation of which node to connect. For example, `['m',0x03]` connects the A/D to node 3.

`rcvBuf=['s',f0,f1,f2]` configures the DAC to generate a cosine wave of frequency $(f2 \ll 16) + (f1 \ll 8) + f0$ and configures the ADC to sample between 100 and 250 times a cycle when possible, with a maximum sample rate of 500Khz. For example, `['s',0xe8,0x03,0x00]` configures the DAC to generate a 1KHz cosine wave and configures the ADC to sample at 250Khz.

5.2 ADC

The ADCS7476 is a successive approximation ADC with an SPI interface. In operation, a master device initiates a conversion by pulling `CS` low, then clocks in data using the `CK` line and reading `SDATA`.

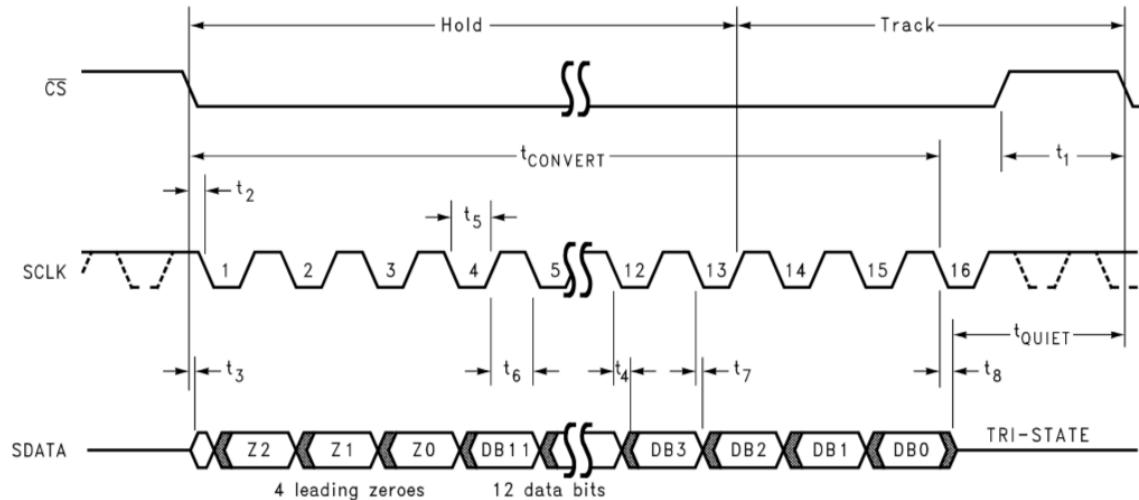


Figure 5-1: ADCS7476 Serial Interface Timing Diagram

The ADC firmware configures two 32-bit timers and the DMA to record data from up to sixteen

external ADCs. The two timers are responsible for driving the Chip Select and Clock lines on all of the ADCs. The timing of the Chip Select line determines the sampling frequency and the Clock line is driven at 8.4MHz. On the rising edge of each clock cycle, the DMA is triggered to sample the data lines of each ADC. A trigger causes the DMA to store each of the 16 pin states on PORT E as a 16-bit integer in the buffer **datas**.

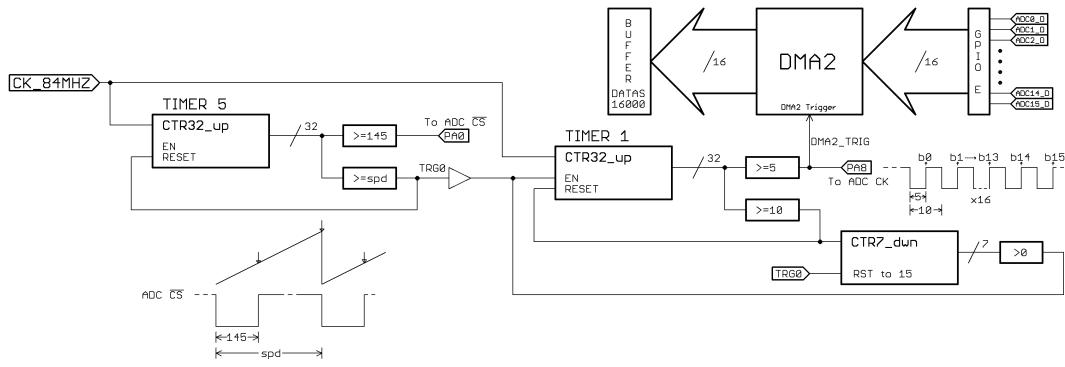


Figure 5-2: Full ADC Driver with DMA

5.2.1 ADC Timers

Timer 5 drives the Chip Select lines. It is configured as a 32-bit up-counting timer clocked at 84MHz, uses output compare unit 1 to control pin PA0 (connected to ADC \overline{CS}), and is set to run in continuous mode. The ADC sampling frequency is controlled by the \overline{CS} line, which is controlled by the period of Timer 5. When a host computer requests an ADC trigger, it sends a sampling frequency along with it. The appropriate \overline{CS} period is calculated and assigned to the variable `spd` (sampling period). Timer 5's output compare mode 1 is set to PWM2 mode, where PA0 is low when the counter is less than the compare value and high when the counter is greater than the compare value. At update, Timer 5's count value is reset to 0, so PA0 falls to 0V and initiates a conversion. When Timer 5 counts to 145, $1.7\mu S$ later, the conversion is complete and PA0 rises to 3.3V, until the count reaches `spd` and the cycle starts again. Timer 5 is also configured in Master Mode and sends a trigger signal on TRG0 at each update event. In this case, the signal on TRG0 is used to enable Timer 1.

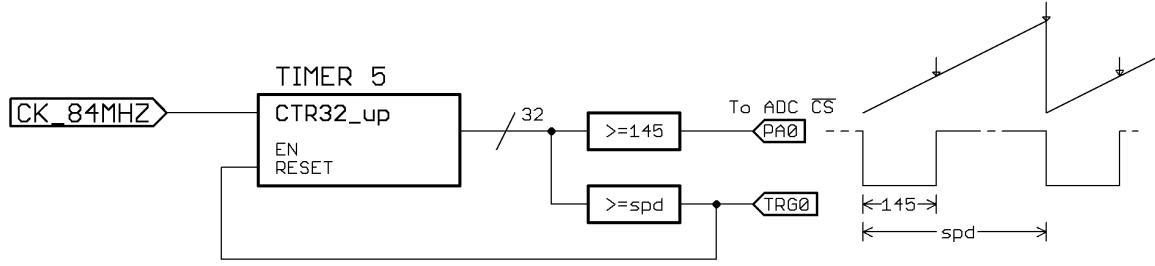


Figure 5-3: ADC Chip Select Timer

Timer 1 drives the ADC Clock lines and behaves as a $59nS$ on - $59nS$ off 16-shot timer. It's also configured as a 32-bit up-counting timer clocked at 84MHz, but it's only enabled by the TRG0 signal from Timer 5. Like Timer 5, it uses output compare unit 1 to control an output pin, in this case PA8. Unlike Timer 5, it's set to run in one-shot mode with a repetition counter that is preloaded with a value of 15. The output compare unit is set to run in PWM2 mode with fixed compare and update registers. When Timer 1 counts to 5, the compare unit fires and PA8 goes high. When the count reaches 10, the timer is updated, PA8 goes low, and the repetition counter decrements. The timer continues to run, toggling PA8 and decrementing the repetition counter, until the repetition counter reaches a value of 0. Once the repetition counter reaches 0, the next Timer 1 update disables the timer.

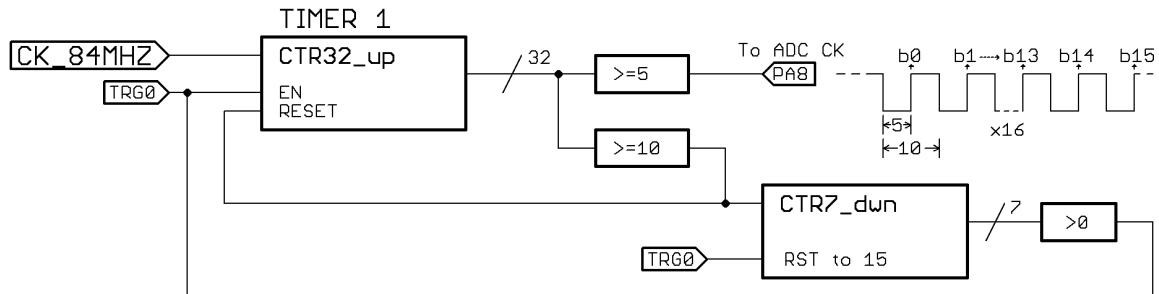


Figure 5-4: ADC Clock Timer

5.2.2 ADC DMA

A DMA controller is used to store serial ADC data from up to 16 ADCs in parallel. On each clock cycle, DMA2 takes the 16-bit integer value represented by the state of the 16 pins on PORT E and stores it in memory. As shown in Figure 5.2, the ADC data pin is updated on a falling clock signal. To record the state of the data pin, DMA2 is triggered to start a data transfer on every compare event from Timer 1, which corresponds with the rising edge of the clock signal.

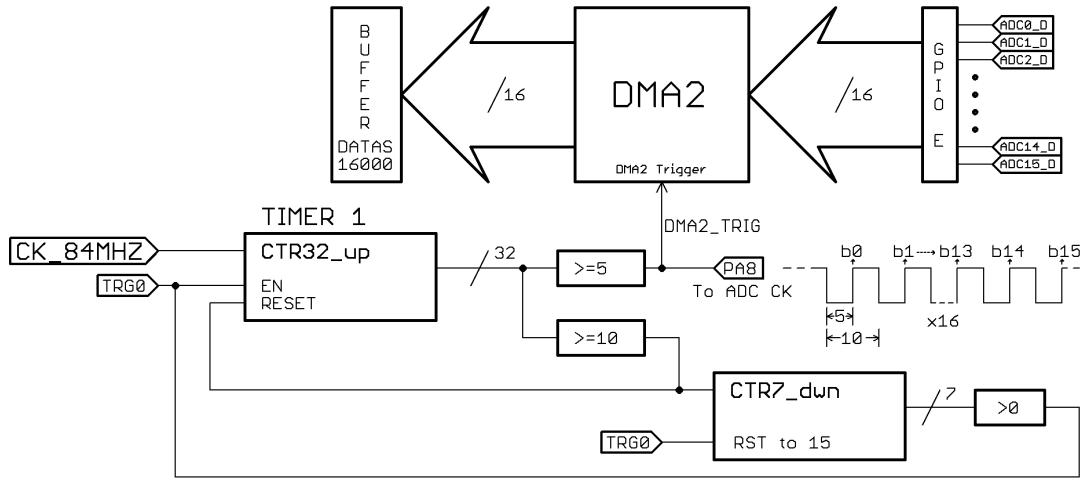


Figure 5-5: ADC Clock Timer Triggers DMA2

When Timer 5 is enabled, DMA2's 'number of data' [to transfer] register is set to 16000 and the ADCs continuously sample at the sample period defined by `spd`. In each cycle of Timer 5, Timer 1 toggles the clock line 16 times, triggering DMA2 on each rising edge. With each DMA trigger, DMA2's number of data register decrements. When the register reaches 0, an interrupt fires that disables Timer 5 and begins the data reconstruction process.

5.2.3 Data Reconstruction

Since the serial data stream from each ADC is stored sequentially, the actual A/D readings need to be reconstructed after they are recorded. In Figure 5.1, the binary data stored in `datas` is examined. Each 16-bit integer in `datas` contains one bit of data from each of the 16 pins on PORT E, which connect to the data pins on the external A/D converters. Here, each letter represents the data from a particular external ADC.

To reconstruct the data collected from $\text{ADC}_a[0]$, a routine iterates through `datas[0:15]`, summing appropriately bit-shifted b_0 's.

$$\text{ADC}_a[0] = (a_{15} \ll 15) + (a_{14} \ll 14) + (a_{13} \ll 13) + \dots + (a_2 \ll 2) + (a_1 \ll 1) + a_0$$

To reconstruct all of the waveform recorded, the routine iterates through `datas[0:16000]`, reconstructing 1000 samples from each ADC.

After reconstruction, the data is sent to a host program over USB for further processing.

Bit Index	b_{15}	b_{14}	b_{13}	b_{12}	...	b_3	b_2	b_1	b_0
<code>datas[0]</code>	p_{15}	o_{15}	n_{15}	m_{15}	...	d_{15}	c_{15}	b_{15}	a_{15}
<code>datas[1]</code>	p_{14}	o_{14}	n_{14}	m_{14}	...	d_{14}	c_{14}	b_{14}	a_{14}
<code>datas[2]</code>	p_{13}	o_{13}	n_{13}	m_{13}	...	d_{13}	c_{13}	b_{13}	a_{13}
<code>datas[3]</code>	p_{12}	o_{12}	n_{12}	m_{12}	...	d_{12}	c_{12}	b_{12}	a_{12}
:	:	:	:	:	...	:	:	:	:
<code>datas[13]</code>	p_2	o_2	n_2	m_2	...	d_2	c_2	b_2	a_2
<code>datas[14]</code>	p_1	o_1	n_1	m_1	...	d_1	c_1	b_1	a_1
<code>datas[15]</code>	p_0	o_0	n_0	m_0	...	d_0	c_0	b_0	a_0
<code>datas[16]</code>	p_{15}	o_{15}	n_{15}	m_{15}	...	d_{15}	c_{15}	b_{15}	a_{15}

Table 5.1: `datas[0:16]`

5.3 DAC

The STM32F407's onboard DAC is used as a signal source for impedance testing. The DAC is configured to update on a timer, and is fed new waveform values from a wavetable by DMA1.

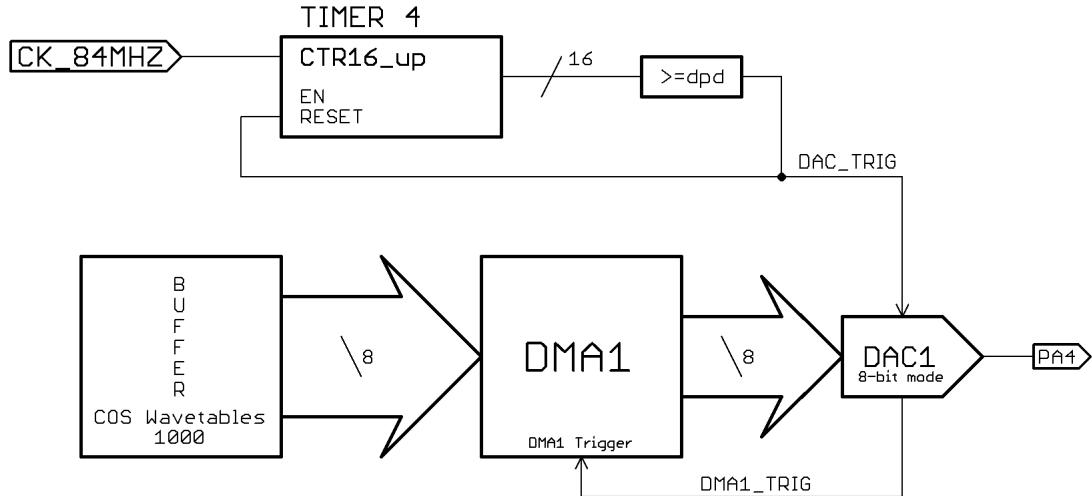


Figure 5-6: DAC Timer Triggers DMA1

Timer 4 is configured as a 16-bit up-counting timer with period `dpd`, set by command over USB. On update, Timer 4 is a master to DAC1, and triggers the DAC to output a new value. The DAC, in turn, triggers a data transfer from DMA1, which transfers the next value in the cosine wavetable to DAC1.

5.3.1 DAC Wavetables

There are three DAC wavetables, each 1000 bytes long. The first wavetable has four cycles of cosine, the second has 10, and the last has 40. The DAC synthesizes sinusoids by taking in a sequence of numbers from a wavetable and converting them, sequentially, to analog voltages. There are two factors that contribute to the output sinusoid's frequency: the rate at which the DAC takes in a new number and the number of samples in the wavetable it takes to complete one cycle. In the case of the first wavetable, where there are four cycles in 1000 samples, the DAC can synthesize a 1Hz signal by converting one number every 4mS, an 0.4% of the period time-step. Had the DAC been looking up the second wavetable, where there are ten cycles in 1000 samples, the DAC would have needed to wait 10mS between samples, a 1% of the period time-step. Using the third wavetable, the wait would be 40mS between samples, a 4% of the period time-step. This amount of delay in the DAC output creates an unacceptable amount of distortion in the output signal.

Wavetables with a higher number of samples per cycle can feed lower-frequency signals at higher sample rates. Additionally, wavetables with a high number of samples per cycle have smaller jumps between samples, resulting in a smoother waveform. According to these observations, a wavetable with an extremely large number of samples per cycle should be used to generate waveforms with the smallest amount of distortion.

However, the DAC's speed is limited by the internal system clock, and to generate sinusoids at 100KHz+, the DAC needs to be fed by a waveform with fewer samples per cycle.

Chapter 6

Software Control

In this chapter, software control of the hardware system is covered. The software that controls the hardware system is very similar to the software simulation in Chapter 3. A control routine parses first circuit 'experiment' requests and configures the hardware to follow out the experiment. The routine then analyzes the returned measurement data. Unlike in the simulation, the measurement data is returned in the time domain rather than the frequency domain. The software transforms the returned data to the frequency domain to calculate the magnitude of the returned waveforms. With magnitude data, the same network sensing algorithm from Chapter 2 is used to reconstruct the network.

6.1 Control Routine

"Just a teaspoon of sugar makes the medicine go down" - Mary Poppins

The control routine is quite self-explanatory.

```
def medicine(drive,chan,freq=1000,adc=0,uglist=[]):
    nog=0          #a number representing nodes not to ground
    for x in uglist:
        nog+=(1<<x)
    ser.write(['g',0xff&(~((1<<chan)|(1<<drive)|nog))]
              #ground all of the nodes whose bits are set to 1
    ser.write(['w',drive+1])
              #connect the voltage source to node 'drive'
    freq = int(freq)
    f=[freq&0xff,(freq&0xff00)>>8,(freq&0xff0000)>>16]
              #send three bytes representing the desired frequency
    ser.write(['m',chan])
```

```

#set the voltage multiplexer to the channel of interest, chan
ser.flushInput()

#clear everything already in the input buffer
ser.write(([‘s’]+f+[adc]))

#tell the microcontroller to begin sampling
# adc value of 0 is for voltage and 1 for current

a = ser.read(2000)

b = [5*((ord(a[2*x+1])<<8)+ord(a[2*x]))/1024. for x in range(len(a)/2)]
b = b[2:]

# the ADC data is 10-bits per sample, but the serial channel only supports
# sending 8-bits at a time, so it has to be reconstructed.

ft=fft(b)

aind=argmax(abs(ft[1:]))+1

#calculate the magnitude of the most prominent sinusoid and return it

amp = abs(ft[aind])/500

return amp

```

6.1.1 FFT

The normalized FFT is used to extract the amplitude data out of the measured sinusoidal voltages and currents. Since the measured voltages and currents should be close to purely sinusoidal, the magnitude of the FFT of those voltage and currents should be close to a delta function at the frequency of interest. The magnitude peak value of the normalized FFT is the amplitude of the measured sinusoidal voltages and currents.

6.2 Network Sensing Algorithm

The implementation of the network sensing algorithm in the hardware control software is identical to that in the hardware simulator, mentioned in Chapter 3.

Chapter 7

Results

In this chapter, results for the simulation and hardware implementations of the RLC network identifying system are presented and analyzed. The shortcomings of the systems are addressed and potential improvements are discussed as future work.

7.1 Simulation Performance

The performance of the RLC network identification system simulator is analyzed by its output precision and cases in which it fails. The output precision of the simulator is quite good and can be controlled almost arbitrarily. The precision depends almost entirely on the thresholds defined during the finite-differencing step, as looser thresholds allow bad data to get averaged into the good data. The cases in which the simulator fails are always cases where a relatively high resistance is in parallel with an LC network. In these cases, the simulator reports that there is no resistor in parallel with the LC network. The failure mode lies within finite difference stencil thresholding, where the simulator never finds a region where the slope of the impedance vs frequency plot is zero, and in turn the simulator decides there is no resistor. Designing a better element identifying algorithm would make excellent future work.

The schematic renderer used to display the returned matrices as a legible schematic could use some work as well. It is effective at creating schematics for circuits composed of resistors with only a few nodes, but the value labeling system and lack of organization quickly gets out of hands when the number of nodes is greater than five. There has been a fair bit of research into auto-generating schematics from netlists, particularly from the days of discrete digital logic, such as *ASG [Automatic Schematic Generator]* from Lageweg at Delft University [?] and the *ASG* from Jehng, Chen, and Parng at National Taiwan University [?].

7.2 Hardware Performance

The hardware RLC network identification system was tested to determine its dynamic range of operation, precision, and ability to detect RLC networks.

7.2.1 Element Testing

Resistance

A selection of resistors ranging from 33Ω to $330k\Omega$ were analyzed with the RLC network identification system. For each resistor, the identified resistance was recorded and plotted.

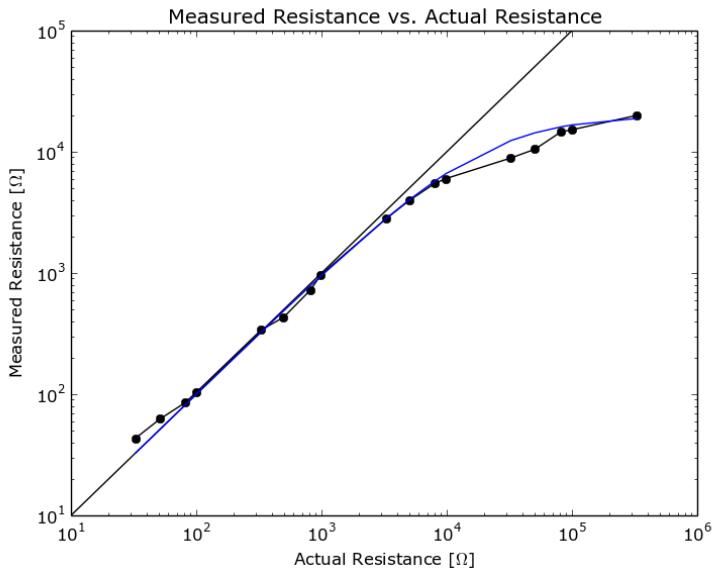


Figure 7-1: Plot of measured resistance for known resistors

From the Figure 7.2.1, the range of resistances for which the system maintains valid precision operation is between ~ 100 and 2000Ω . For resistances larger than $2k\Omega$, the system returns resistances that are far smaller than the input resistance. The asymptotic portion of the curve in Figure 7.2.1 indicates a parallel resistance somewhere around $20k\Omega$. The blue line indicates the equivalent resistance of each tested resistance with a parallel $20k\Omega$ resistor. This parallel resistance is not real. Using an ohmmeter, it was verified that there is not enough leakage current between breadboard nodes or between breadboard nodes and ground to account for $20k\Omega$ of resistance. Rather, this measurement error is caused by the minimum non-zero signal amplitude measured on the output of the current amplifier A/D converter. A 10-bit ADC that samples a $5V$ window converts every $.004V$ into one bit. The minimum amplitude signal produced by a 10-bit ADC sampling over a $5V$

scale has an amplitude of $4mV$. The test voltages applied to each node have amplitudes around $0.75V$. The current amplifier has a gain of 100, so the effective current amplitude measured is $40\mu A$. $0.75V/40\mu A = 18750\Omega$

So, the maximum possible impedance measurement from any node to ground, $R_{n||}$, is $19k\Omega$. A frequency sweep of the impedance of an open node, such as the one in Figure 7.2.1, demonstrates this.

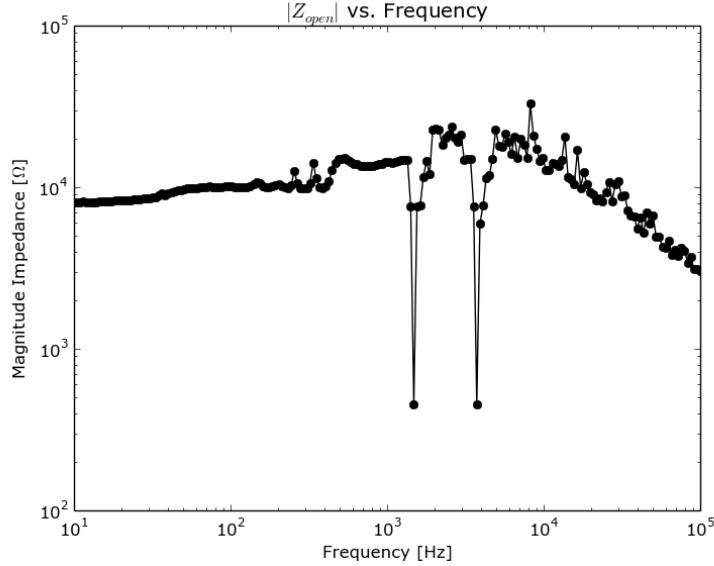


Figure 7-2: Driving point impedance of an open node from 1Hz to 100KHz

By adding a variable gain stage to the current amplifier, or a variable-sized sense resistor, the maximum impedance measurement could be modified.

The maximum impedance measurement between two nodes, rather than just one node to ground, is an order of magnitude higher. This is due to the procedure used in the network sensing algorithm, where $V_t R_{n||}$ is divided by V_m . If V_n is smaller than 1, the resulting resistance R_{nm} is larger than $R_{n||}$.

On the other side of Figure 7.2.1, for resistances smaller than 100Ω the system returns resistances that are far larger than the input resistance. This can be attributed to the DAC buffer current limit, which is $\pm 25mA$. For a test voltage signal of $\sim 1.5V$, a 60Ω resistance produces a $25mA$ current. This agrees with the recorded data quite well. To measure resistances smaller than 100Ω the output buffer could be replaced with one that has a higher output current limit. Alternatively, a variable-sized sense resistor would also mitigate the problem.

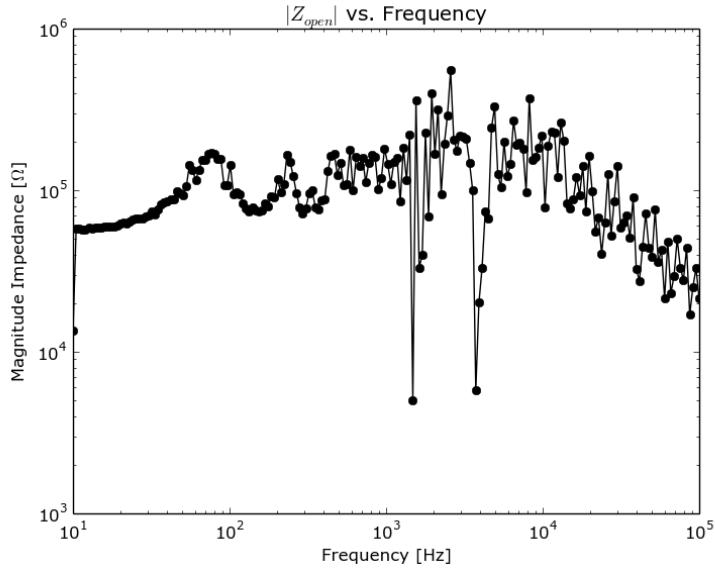


Figure 7-3: Impedance between two open nodes from 1Hz to 100KHz

Capacitance

In addition to the resistors, a selection of capacitors ranging from 220pF to $22\mu\text{F}$ were analyzed with the RLC network identification system. For each capacitor, the identified capacitance was recorded and plotted in Figure 7.2.1.

The RLC network identification system can precisely measure capacitances over three orders of magnitude, from 2.2nF to $10\mu\text{F}$. On the low end of Figure 7.2.1, the parallel parasitic capacitance is visible, which adds about 150pF to the system, although this does not entirely account for the measured capacitance. On the high end, the impedance of the capacitors at the test frequencies drops below 50Ω , again causing the buffer op-amp to reach its output current limits and saturate. The working range of 1nF to $10\mu\text{F}$ is acceptable for general purposes.

Inductance

A selection of inductors were analyzed with the RLC network identification system to no avail. The identification system was not able to find the inductors due to the noisy nature of the impedance measurements that inductors were found to return.

The 'noisy' impedance, shown in Figure 7.2.1, defeats the finite-differencing technique, and causes it to return a large (amplitude of 20) random-looking signal. This is much larger than the +1 output that is expected for an inductor, and causes the element identification step to fail.

In addition to recording noisy measurements while attempting to identify an inductor alone, noisy measurements were found when attempting to identify RL, LC, and RLC parallel circuits. One

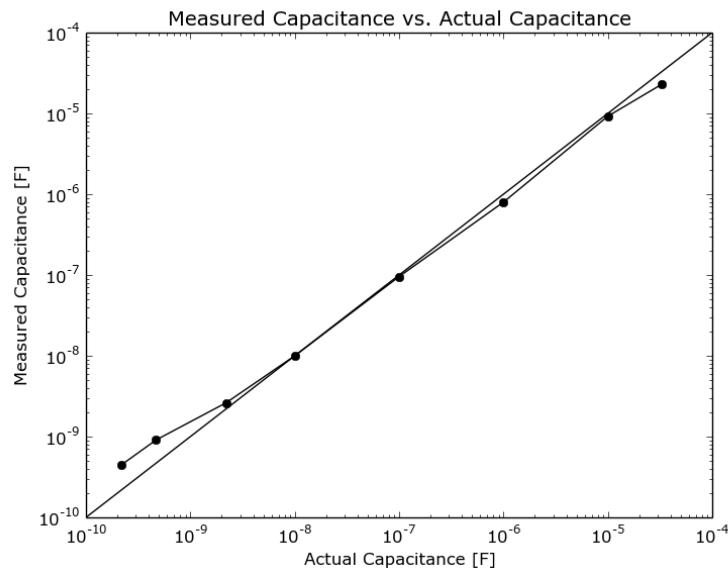


Figure 7-4: Plot of measured capacitance for known capacitors

would expect to find noisy inductive data in the section of the plot where the inductor dominates and clean capacitive or resistive data in the sections of the plot where the capacitor or resistor dominates, but this is not the case. This issue has not been closed, yet is integral to the acceptable operation of an RLC network identification system. Identifying and solving this problem would make excellent future work.

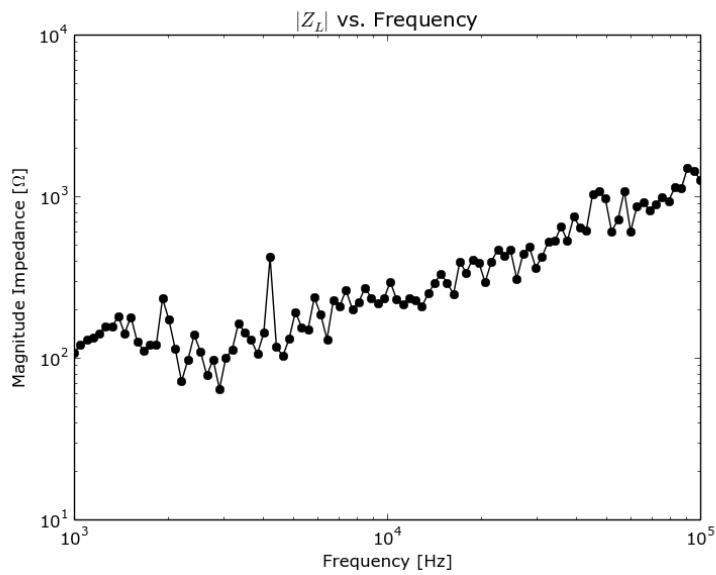


Figure 7-5: Impedance vs frequency for a $3.3mH$ inductor, demonstrating inductive behavior albeit noisily

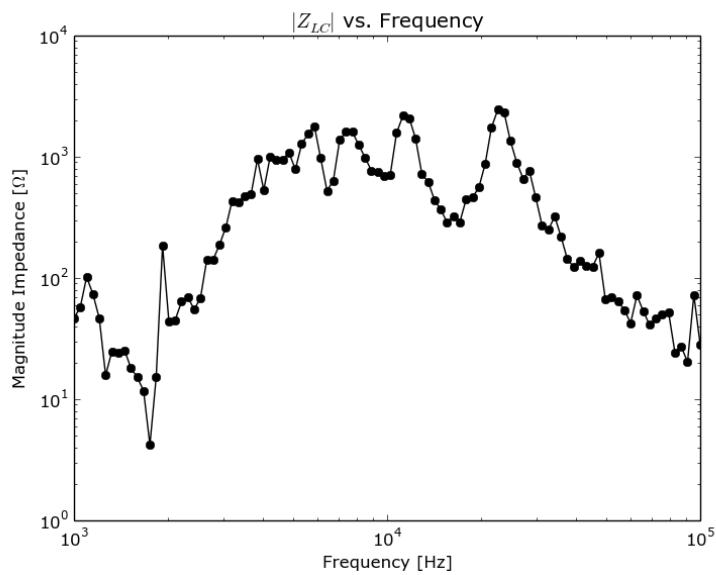


Figure 7-6: Impedance vs frequency for a $3.3mH$ inductor in parallel with a $0.15\mu F$ capacitor, demonstrating LC behavior albeit noisily

7.2.2 Network Testing

Resistive Networks

To verify the network identification system performed properly, complete networks of three, four, five, and six nodes were constructed from resistors within the range of 100Ω to 1000Ω .

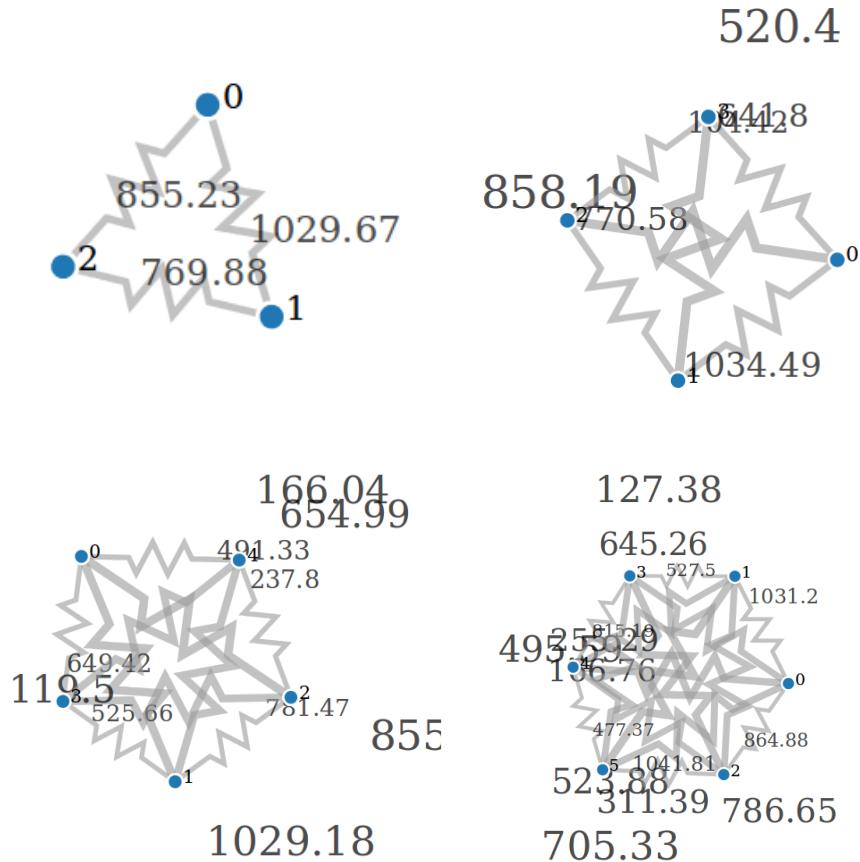


Figure 7-7: Schematics of 3, 4, 5, and 6-node resistor networks rendered in the browser

The RLC network identification system successfully detected all of the resistive connections in the complete 3, 4, 5, and 6 node networks. However, many of the values were detected with more than 5% error once the network became large. This error is produced by the nonlinear weighting that occurs when combining resistors in parallel. As the network grows, each $R_{n||}$ measurement contains more resistive elements in parallel with the branch of interest. If all of these resistive elements are on the same order of magnitude in value (R , the additional attenuation due to the addition of another parallel resistor goes as $\frac{k}{k+1}$, if k is the current number of resistors. This causes error in the measurement of R_{nm} to be amplified by $R(k + 1)$

$$\begin{aligned}
R'_{nm} &= \frac{V_t R_{n||}}{V_n} \quad \text{from Section 2.1.3} \\
&= \frac{\frac{r_{nm}}{R+r_{nm}}}{\frac{1}{R+R_{nm}}}
\end{aligned}$$

with error Δ_r, Δ_v :

$$\begin{aligned}
R'_{nm} &= \frac{V_t R_{n||}}{V_n} = \frac{\frac{R_{nm}}{R+kR_{nm}} + \Delta_r}{\frac{1}{R+kR_{nm}} + \Delta_v} \\
&= \frac{R_{nm} + \Delta_r(R + kR_{nm})}{1 + \Delta_v(R + kR_{nm})}
\end{aligned}$$

If $R_{nm} \approx R$, error contributed by $R_{n||}$ measurements is $\Delta_r R(k+1)$ and the error contributed by V_n measurements is $\frac{1}{1+\Delta_v R(k+1)}$, both of which are scaled by the number of parallel resistances. To improve performance in the cases where many branches connect to one node, improve measurement precision and dynamic range. In the previous tests, the average resistance was $\sim 500\Omega$. Five 500Ω resistors' equivalent parallel resistance is 100Ω , close to the lower bound on the identification system's dynamic range. This analysis explains the asymmetry found in many of the measured resistance matrices, since each node had different parallel equivalent resistances, $R_{n||}$.

The following data correlates to the test networks in Figure 7.2.2. Column index i and row index j are node numbers, the value at index i, j is the resistance R_{ij} .

N=3

R measured			R actual		
None	1.03e+03	8.55e+02		1.00e+03	8.20e+02
1.03e+03	None	7.70e+02		1.00e+03	7.50e+02
8.56e+02	7.71e+02	None		8.20e+02	7.50e+02

N=4

R measured			R actual		
None	1.03e+03	8.58e+02	6.42e+02		1.00e+03
1.03e+03	None	7.71e+02	5.20e+02		7.50e+02
8.63e+02	7.84e+02	None	1.04e+02		8.50e+02
6.46e+02	5.27e+02	1.04e+02	None		6.20e+02

N=5

R measured	R actual
None 1.03e+03 8.56e+02 6.49e+02 4.91e+02	1.00e+03 8.50e+02 6.20e+02 5.10e+02
1.04e+03 None 7.81e+02 5.26e+02 1.66e+02	1.00e+03 None 7.50e+02 5.10e+02 1.60e+02
9.99e+02 8.87e+02 None 1.19e+02 2.38e+02	8.50e+02 7.50e+02 1.00e+02 2.00e+02
6.66e+02 5.31e+02 1.05e+02 None 6.55e+02	6.20e+02 5.10e+02 1.00e+02 6.20e+02
5.02e+02 1.67e+02 2.10e+02 6.54e+02 None	5.10e+02 1.60e+02 2.00e+02 6.20e+02

N=6

R
None 1.03e+03 8.65e+02 6.45e+02 4.96e+02 5.24e+02
1.04e+03 None 7.87e+02 5.28e+02 1.67e+02 7.05e+02
1.06e+03 9.65e+02 None 1.27e+02 2.55e+02 1.04e+03
8.10e+02 6.60e+02 1.31e+02 None 8.15e+02 3.11e+02
5.57e+02 1.91e+02 2.39e+02 7.67e+02 None 4.77e+02
5.19e+02 6.97e+02 8.51e+02 2.49e+02 4.12e+02 None

RC Networks

Similar results are found when measuring RC networks.

7.3 Improvements

Various improvements would be welcome future work. From design focused on individual subsystems to topological changes, there is much to improve on this iteration of a circuit sensing breadboard.

7.3.1 Variable Gain Amplifiers

Adding a Variable Gain Amplifier to the ADC, DAC, and current sense amplifier signal chains would enable a large increase in dynamic range of the entire system. Not only would it improve the resolution (and thus precision) of measuring elements on the far-ends of the dynamic range, but it would enable higher resolution captures of small signals, improving the overall utility of the device.

7.3.2 Variable Sense Resistor

Similar to the aforementioned improvement, this would improve the dynamic range and resolution of the current measurements.

7.3.3 Better Switches

The switches in this system are the main contributors the high parasitic capacitance between every node and ground. The parasitic capacitance causes undesired attenuation of signals at high-frequencies, making This version of the circuit sensing breadboard unsuitable for even moderate-frequency circuit design. Finding switches with lower parasitic capacitances would likely require permitting higher $r_{DS_{on}}$, which is possible but requires further testing. Alternative technologies for switches, such as relays, seem promising as they have very low parasitic capacitance and low on resistance, at the cost of bulk.

7.3.4 Better Schematic Display

The schematic display was built on top of D3, a javascript library for graphical applications in the browser. D3 was chosen because of its large community and available relevant example code. In particular, *Force Graph*. Force graph is interesting because it could be configured to take an arbitrary graph and simulate a kind of spring fabric between the nodes. Nodes tied together with edges are attracted to one another, but nodes not connected repel. This action 'unfolds' random networks until they appear flat, with as few crossing branches as possible. Replace the branches with resistors, capacitors, and inductors and voilia! Unfortunately, using a relatively unmodified version of force graph in this application makes schematic display design lacking in many facets. Namely, lack of well-placed labels, trouble with handling parallel components, lack of orientation, and lack of visual control. Redesigning the current schematic display or designing a new one would greatly improve the usability of the system.

7.3.5 Additional Devices

At the moment, the RLC network identification system can only identify resistor and capacitor networks with a maximum of eight nodes. Designing a system that can successfully identify inductors would be a huge improvement. Additionally, determining how to identify systems composed of more complex components like diodes, transistors, transformers, and even integrated circuits is an interesting and likely rewarding problem to solve.

7.3.6 Computer-Aided Debugger

Although the intent of this hardware system is to display a live schematic of a circuit under construction, the hardware is capable of serving other purposes. Take, for example, a system that is given the schematic of the circuit at hand. The circuit is performing unexpectedly, so the circuit symptom is reported to the system. The system, knowing what the circuit should be, computes how the circuit should behave. Now, since the system knows how the circuit should behave, it asks

the user to probe various parts of the circuit, comparing its simulation of the circuit to the actual signals on the circuit. Mismatches in simulated and physical signals are traced to their sources, and the bug is found. If the system were also give electrical access to every node in the circuit, it may also be able to debug the circuit on its own. This would make excellent future work.

7.4 Conclusion

A system that draws a live schematic of a resistive circuit built on an eight-rail breadboard was presented and analyzed. The system was decomposed into its algorithmic, hardware, firmware, and software components, each of which were described in detail. A new method for mounting breadboards to printed circuit boards was prototyped and presented. The performance of the overall system was measured and analyzed, and much was learned along the way. Although RLC network identification systems may not be feasible, realistic, or desirable, future work on this topic could lead to some interesting and novel systems.